

User Behaviour in Gestural-Interface Environments

Martin Henschke

A thesis submitted for the degree of
Doctorate of Philosophy at
The Australian National University

November 2016

Draft Copy – 25 November 2016

© Martin Henschke 2016

Except where otherwise indicated, this thesis is my own original work.

Martin Henschke

25 November 2016

Dedicated to my mother and father, who led by example to become a researcher

Acknowledgments

I wish to acknowledge the role my supervisors had in assisting me in penning this thesis; my primary supervisor Prof. Tom Gedeon for his research expertise, guidance and instruction that ensured both my research and my publications were successful. My co-supervisor Richard Jones was invaluable in the preparation of my work and managing the project.

I also wish to acknowledge other members in my the iHCC research group and my peer group, especially Sabrina Caldwell for providing direction in my research and acting both as a participant in my experiments and proof-reader of much of my work.

Finally, I acknowledge and thank the many friends I've made at ANU for their past and continued support, and for making the years I spent as a PhD student as enjoyable as they were.

Abstract

Put your abstract here.

Contents

Acknowledgments	vii
Abstract	ix
1 Introduction	1
1.1 Topic	1
1.2 Research Question	2
1.3 Thesis Outline	3
1.4 Significance	4
2 Natural Language and Gestures	5
2.1 Inter-communicative gestures	5
2.2 Forms of inter-communicative gestures	6
2.3 Elements of gestural communication	7
2.4 Cultural variation and significance	7
2.5 Variation in performance of gestures	8
3 A Survey of Gestural Interfaces	11
3.1 Definition	11
3.2 Intended Purpose and Origins	11
3.2.1 Natural Language Interaction	11
3.2.2 Expressiveness	12
3.3 Development and Implementation	12
3.3.1 Two-dimensional / iconographic	12
3.3.2 Pen and tablet input	13
3.3.3 Motion sensing and towards "hands free"	13
3.3.4 Gesture Capture Devices	13
3.4 Techniques in Recognition	13

4	Implementation and Usage of Gestural Interfaces	15
4.1	Pointing Interfaces	15
4.1.1	Implementation	15
4.1.2	Design	18
4.1.3	Implementation Methodology	20
4.1.3.1	Design Constraints	21
4.1.3.2	Pointing Recognition	24
4.1.3.3	Practical Issues	28
4.2	System Performance	30
5	Performance with Pointing Gestures	31
5.1	Actual and Perceived Interface Functionality	31
5.1.1	Natural Computing	35
5.2	Study 1	37
5.3	Changing the user perception	39
5.4	A Comparison of Pointing Devices	41
5.4.1	Aim	41
5.4.1.1	Performance	42
5.4.1.2	Ergonomics	42
5.4.1.3	Preference	43
5.4.2	System	43
5.4.2.1	Wired Glove	43
5.4.2.2	Wand	44
5.4.2.3	Gyroscopic Mouse	45
5.4.3	Method	46
5.4.3.1	System	46
5.4.3.2	Tasks	47
5.4.3.3	Measurements	49
5.4.4	Results	49
5.4.4.1	Performance	50
5.4.4.2	Ergonomics	51
5.4.5	Discussion of Results	53

5.4.5.1	Trial Caveats	54
5.5	Followup Studies	55
5.5.1	Changes in Study Method	56
5.5.2	Standing and seated postures	58
5.5.2.1	Results	59
5.5.2.2	Discussion	60
5.5.3	Hand shape	61
5.5.3.1	Results	62
5.5.3.2	Discussion	62
5.5.4	Identifying Optimal Attributes for 3D Interfaces	63
5.5.4.1	Introduction	63
5.5.4.2	Methodology	64
	System	65
	Input Devices	66
5.5.4.3	Experimental Methodology	67
5.5.4.4	Results	68
	Performance	68
	Preference and Fatigue	69
	Discussion	70
6	Performing with Freeform Gestures	73
6.1	Future Work	73
7	Conclusion	75
7.1	Who even knows...	75

List of Figures

List of Tables

5.1	Mean and standard Deviations for time and errors in completing selection tasks for each device	50
5.2	Distribution-free analysis of questionnaire results pertaining to immediate and continually-induced fatigue from each device	52
5.3	Difference in normalized arm angle between standing and sitting trials	60
5.4	Difference in normalized arm angle between standing and sitting trials	60
5.5	Difference in normalized arm angles between free hand and splinted hand	62
5.6	Mean Times and Standard Deviation for each Device in Performing Selection Trials	68
5.7	Mean Times and Standard Deviation for Each Device in Performing Grab and Drop Trials	69
5.8	Mean change of average dynamometer recordings from baseline to after trial with device	70
5.9	Reported fatigue from each interface on a scale from 1 to 7	70

Introduction

The introduction will define the topic of the thesis. It's goal is to provide context to the reader for the information presented in the rest of the body of the work. It is non-technical, gestalt and comprehensive.

1.1 Topic

The topic provides an initial discussion into what the thesis is about on the whole. It provides the first introduction to the topic from a simplified perspective, and how the thesis addresses the topic matter.

This thesis explores the topic of interface that makes use of human hand and body movements as a modality of input, referred to as gestural interfaces. Such interfaces, which measure this movement either through the use of sensors attached to the body or through cameras, allow humans to interact with interfaces either in a way that better represents the task they are trying to perform, such as kicking a ball in a virtual environment, pushing a 3D object or pointing to an area of the screen to interact with directly. This perception of ease of use combined with the opportunity for more expressive input and novelty factor have seen gestural interfaces becoming commonplace consumer devices.

However, with this increase in consumer interests, more and more attention is drawn to the drawbacks of touchless gestural interfaces. Limitations in the technology in terms of gesture capture and classification mean such interfaces can be difficult for individuals to use; gestures are typically defined by a very rigid requirement for

movement which can be difficult to replicate, especially if the movement required is not communicated sufficiently. Even then, performance of such interfaces is never 100% and given the smaller set of interactions performed and the higher time and effort investment required to perform them, this lower recognition rate can detract from the user experience negatively. Ergonomics are also an increasingly prevalent concern; most computer input requires low physical movement and effort to operate, and the higher requirement for physical movement can fatigue users. Poorly defined gestures that demand repetitive movement, hyperflexion or unnatural joint movement can also lead to strain and discomfort even in short sessions of use.

These issues are sufficiently prevalent that work into finding more ergonomic, comfortable and useful gestures for typical computer usage has become the predominant focus of many such technologies in the past few years. This thesis intends to add to the body of knowledge on the subject by exploring factors specific to how humans interact with such interfaces; specifically treating such interfaces as natural, or inherent in their own behaviour and observing how they would interact in a low-constraints environment. The ultimate goal of this research is to find a set of guidelines or constraints to match as many users as possible, or find where flexibility is necessary in gestural interfaces, with the ultimate goal of making them both more comfortable and more intuitive to use in future iterations.

1.2 Research Question

This discusses in detail the questions to be answered in this thesis

We wish to answer the following questions in this research:

1. Can a user interact with a gestural interface naturally? When given a low-constraints system, the natural interface hypothesizes there is a natural method in which users will wish to interact given how the display is provided and the obvious system constraints made visible, such as environmental features (controllers, or how the display presents the information shown). If given a loosely-constrained system, what gestural interactions would a user naturally use? Are there limitations or

requirements that enforce this?

2. How do users reconcile fatigue with interaction? Such interfaces typically produce significant fatigue. How would the user adjust their performance to avoid this, if at all? How can we make our systems reactive of this?

3. How do users prefer to interact with gestural systems? Do they prefer a highly-constrained environment or a more relaxed environment? What kind of gestures do they most like to perform when doing typical tasks? Do we see abstraction or highly iconic gestures being preferred for most interactions?

1.3 Thesis Outline

The manner in which the research was conducted and how it is presented is discussed in this section.

Chapter 2 of this text describes gestures as a vector of human-human communication, and some of these applications in the realm of human-computer interaction. This includes the structure and syntax of natural gestures, their purpose in communication, the similarities and differences in cultures and groups. It also touches on the use of direct, natural languages like ASL and Kodaly notation in computer usage.

Chapter 3 is a review of touchless gestural interfaces. It moves from early gestural interfaces, taking a historical perspective and analysing key features and purposes of systems found in the literature, before systematically comparing systems to quantify function and utility of such systems in contextualizing the use-cases this thesis is exploring.

Chapter 4 discusses the design and implementation of a deictic gestural interface. The context of the system design and purpose is explained, as well as how it was used. Issues with implementation and lessons learned from the design are also covered as background material. This is all presented for discussion in the next chapter.

Chapter 5 describes the experiments run using the gestural control interface. It covers a series of experiments that involve variations on the interface, controlling various

aspects of how the participants interact. Results are presented with discussion.

Chapter 6 describes another set of experiments run exploring users with a freely-defined gestural interface, and variations on that set of experiments. Results are presented as before.

Chapter 7 discusses the results in the context of the research question. It explains their significance and extrapolates on how they can be applied and their value in designing gestural interfaces for user interaction, and concludes the work.

There are two appendices: Appendix A describes the application of the pointing interface to a data visualization, while Appendix B explores their value in human-robotic interfaces.

1.4 Significance

Discusses the impact of the work in the context of the problem and existing research

Potentially a topic for the conclusion rather than introduction-discuss with supervisor

Natural Language and Gestures

The introduction to the chapter should discuss natural gestures as clearly and succinctly as possible.

2.1 Inter-communicative gestures

Introduce inter-communicative gestures as a concept. Define and give examples.
Introduction and motivation for the chapter

Gestures are a form of non-verbal communication that are performed largely subconsciously. They function as a vector of communication independently, or integrated with spoken language. The amount and kind of information gestures can convey depends on the form of gesture, but this can include things from basic context in spoken conversations, clarification or disambiguation of spoken concepts, conveying affective information and attitudes of the communicator, elaborating or demonstrating, particularly in reference to a physical place or task, or conveying information subtly. The majority of inter-communicative gestures are performed by the hands and arms, and often in front of the body or around the face of the converser. Other aspects of indirect gestural communication, such as posture and gait can also affect this information.

This thesis begins with an examination of the natural occurring context of gestures to better understand their role in basic human communication, and how best they would intergrate with human-computer communication.

Sources: Argyle [1988], Knapp and Hall [2010], Cassell [1998].

2.2 Forms of inter-communicative gestures

Discuss the differences between iconic, deictic, metaphoric and beat gestures as features of inter-communicative language, as well as ergotic gestures used in human-computer interaction

Four forms of inter-communicative gesture: Iconic, deictic, metaphoric and beat gestures. These are all used when integrated with human speech, although some can be used externally.

Iconic gestures are also known as 'miming gestures'. These involve performing hand or body movements that are a facsimile or equivalent to the gestures performed in doing a specific task with a given object or in a certain environment. The gesture clarifies to the observer what the action is, or may point out how the action is done, what it involves etc.

Metaphoric gestures are specifically inter-communicative and are used to reference conversational objects or topics as though they had a physical presence. This is mostly used to disambiguate non-declarative language; where the indefinite is used (e.g. 'I was talking to him') a metaphoric gesture may point to a previously referenced physical space to clarify.

Deictic, or pointing gestures serve the purpose of associating a physical object or direction with the topic of discourse to the interlocutor.

Beat gestures are naturally-occurring with spoken language, are broadly universally performed (with variation, see Argyle [1988]) as a means of adding emphasis and stress to words or sections. This allows certain sections that are peripheral or unimportant to be skipped, while the most important points are better emphasized.

Emblematic gestures can be considered a subset of iconic gestures, or may exist on their own; these are gestures that have an innate (typically simple) meaning, typically indicating commands for the interlocutor or very simple attitudes. Sets

of these typically exist in all languages; and though they may have evolved from a specific event or feature of the natural world they are now broadly arbitrary. Complex syntaxes of emblematic gestures can form their own languages; such is the case for code signs used by police and sign language used by the deaf.

2.3 Elements of gestural communication

A discussion of all elements involved in secondary communication. Includes prevalence and significance of using gestures to communicate, as well as the hand and body movements made, or other features (like eye contact, or simpler/subtler movements)

2.4 Cultural variation and significance

Covers a number of references exploring gestural performance and variation between different cultural groups, including separate countries, sign language communicators, and broader groups. Contrasting both similarities and differences

Gestures are a universal component of language; no living language is known to explicitly exclude their usage, and they are generally thought to provide additional information to spoken language, making it better to understand and easier to follow.

Gestures are observed in non-human species, particularly primates. Many breeds of monkey have been observed manipulating their hands and pose to express different information, particularly affective information about their mood or their relationship to other monkeys (submissive, dominant, sexual etc.). Such expressive gestures can involve hand, head or full body movements, often accompanied with cries or a specifically adopted posture. Although uncommon, monkeys have been observed using pointing gestures, and many animals are capable of understanding humans when pointing. Goodall [1968] Argyle [1988]

Gestural language is considered a universal feature of human communication. It has been hypothesized that gestures may have been one of the earliest modes of com-

munication, developing before or concurrently with speech, given it has been shown breeds of gorilla and chimpanzee are able to learn and use sign language competently without further evolution, while the same is not true for vocal communication. Hewes [1976]. Some forms of gestures can be found in all living languages, and there exists a small subset of gestures that are considered linguistically universal; these are recognized and used in all cultures with a reasonably similar meaning; it is expected such gestures may have evolved from natural behaviour or very early human interaction. Many such gestures however have although the same performance can have completely different interpretations depending on the culture Knapp and Hall [2010]. This can mean that gestures not only have many different meanings depending on the language or region they are performed, but they can have very different meanings; for example the thumbs up, which is a simple acknowledgement in most languages, can be highly derogatory or offensive in others.

It can be expected that all computer users proficient in a natural language will have no difficulty understanding how gestures may be performed to operate a computer system, especially if they are already familiar with computer systems or it is presented in the context of typical human communication (see ?). However, a single set of gestures will not be sufficient for a broad user base, especially if that base is multicultural due to the very high level of variation. Certain gestures may be unintuitive or complicated to the performer as they have no cultural significance, others may find them uncomfortable to perform due to cultural connotations or illogical given the circumstances. A common example is the direct deictic point; although a simple and very common gesture in both spoken language and computer usage, pointing with the index finger at someone can be considered highly rude in some cultures, limiting the utility of this gesture in certain situations without forcing users to behave unnaturally.

2.5 Variation in performance of gestures

Gestures are a part of all languages, though are performed different and mean different things between cultures. The question of whether or not they are performed

differently by different groups of individuals or even if gestural performances remain consistent with an individual marks at the core of what 'comprises' a gesture. Most gestures are defined by the general motion made or the positioning relative to other parts of the body, where as recognition is often more specific. Isolating what exactly people naturally think of in defining gestures, and how that might change the way they perform them over time, would help us to better understand how to adapt to users over many repeated performances, as well as how to define our systems in general.

This will be discussed in more detail in chapter 4.

This section is currently undercited. More results can be found to fill it out.

relabel

This is moving towards the research done in experiment (discussed in chapter 4 in more detail). Potentially worth moving to a later slot.

A Survey of Gestural Interfaces

3.1 Definition

We define gestural interfaces, and their various flavours, types and contexts here.

What is a "gesture" in terms of a computer system?

These are defined according to basics texts on the subject. Wixon and Wigdor [2010].

Hand-based gestures and taxonomy as in ?.

Breakdown of general features, a gesture taxonomy and method of capture by classifiers

State-based, or "static" gestures versus "fluid" gestures, that encapsulate movement

Devices used in their capture, and the sort of gestures each captures

3.2 Intended Purpose and Origins

3.2.1 Natural Language Interaction

Factoring from chapter 2

Gestures are a method of performing interaction with a computer system that mimics or copies the manner in which we communicate with other humans

This opens large possibilities for less direct and more ambiguous input when controlling and interacting with computer systems.

Discussing natural interaction, ?. Use of existing sign languages (ASL, Kodaly), see: LicsÅar et al. [2006], ?, Woodward [1976].

The merge of natural language with interaction.

Spoken language as a modality (particularly Bolt [1980]).

Uses in relation to typical tasks English more prevalent in searching etc. and "matching" gestures to equivalent tasks (iconic representations). Consumer examples include Nintendo Wii, therapy papers on related topics also relevant.

Control of interfaces for motion relevant in diverse environments (see ?, ?).

3.2.2 Expressiveness

Not a lot of material on this subject. To review.

Limits in information transmission through unary inputs

Multi-dimensional addition given by gestures

Mult-modalities with gestural interfaces

Semantic richness; embedding large amounts of meaning into smaller sections of text.

3.3 Development and Implementation

This covers the history and development of the field of gestural interaction and input, exploring it from the earliest examples of body language capture to the modern devices used today.

3.3.1 Two-dimensional / iconographic

Historical perspective. The earliest gestural interfaces did very basic capture, either comparing images or one-dimensional input combined with other information for

very basic input. Exploration should include very early pointing interfaces as handled by proximity sensors (Bolt [1980]), Sketch pen interfaces (citation needed), and the like.

3.3.2 Pen and tablet input

Shorter Section. Looking at basic gestural interfaces that are captured using tablet and pen, and their development. Early examples include Graffiti and optical character recognition interfaces. Move onto quick-type systems and towards the touch-based screens and devices we used today. The most unilateral and present form of gestural interface encountered by modern consumers.

3.3.3 Motion sensing and towards "hands free"

Earliest examples of camera solutions in computer science. Bulk of history can be explored here, due to the rapid rise in interests in the mid 80's and early 90's. Section should include camera efforts from early systems, triangulation, then moving towards depth cameras and combined inputs. See endnote for citation list.

3.3.4 Gesture Capture Devices

Specialized capture devices, that mimic previous efforts. Discuss improvements in accuracy, particularly with the introduction of accelerometry. Includes data gloves, hand-held controllers with sensors etc.

3.4 Techniques in Recognition

NOTE: This is probably the hardest section of the chapter. It's outside my field of confidence and is possibly an unnecessary edition. To discuss with supervisors, but a large body of work from ANU and beyond can be called on. Potentially a reference to CÃ¡l'dras and Shah [1995] is sufficient.

Implementation and Usage of Gestural Interfaces

4.1 Pointing Interfaces

4.1.1 Implementation

We define pointing interfaces in terms of a deictic gesture. A deictic gesture can be defined as any physical movement that indicates to the observer something of relevance within the context of the communication. In the context of a computer system, such gestures are used to indicate to the system an item or area that the user wishes to interact with. This is accomplished by moving a digital 'cursor' to the location of interest before performing an operation such as a mouse click to signal to the system the interaction to be performed at this area.

The analogous deictic gesture used in inter-personal communication is seen in the form of a physical movement to provide a directional or locational indication to the interlocutor [Cassell 1990]. Most unambiguously, an outstretched arm pointing at an object indicates its significance, but a more subtle movement of the finger, head or even other parts of the body can be used in making this indication. In disambiguating the meaning of the gesture, the converser can explain the purpose of the communication ("I want that one", "It's in that direction"), or such context can be gained from earlier conversation.

cumbersome
wording

There are a number of features of deictic gesturing between humans that is lacking

in a WIMP interface in terms of expressiveness and semantic information conveyed. Both draw attention to a place or thing, then indicate the significance of their indication, but there are some degrees of expressiveness physical gesturing affords that a mouse cannot. A physical gesture can indicate additional information in a conversation through the manner of the point, particularly in its obviousness to the observer; an important or rushed gesture can be made more exuberantly. Conversely, a point designed not to draw the attention of other observers or that of a sensitive nature may be made with a subtle head movement or even a break in eye contact to discreetly convey additional information. Pointing can also be accompanied with disambiguating information, in the event a precise point is difficult to perform. For example, if attempting to differentiate a choice on a number shelves, spoken text indicating the item's colour or relative position to clarify the gesture. The point can also be given additional meaning through descriptive conversation; when providing instructions the deictic gesture may only be able to indicate directionality but additional instruction like 'at the end of that street, then to the left' can be used not only to make more sense of the gesture but to allow the gesture further expressiveness, such as pointing in which direction to turn.

Overall in spoken communication, deictic gestures can be given a high degree of expressiveness due to the properties of body articulation and additional commands to which a computer deictic gesture traditionally lacks. A mouse is able to indicate a specific point on a screen to interact with, then have a reasonably finite and very discrete methods of expressing intent. For the majority of computer use, this is perfectly appropriate as ambiguity is not a natural interact with such systems, and most humans operate computers as they would a tool rather than as an intelligent entity they give vague commands to.

This section discusses a 3D pointing interface for a computer system. Such an interface is one that allows users to perform pointing tasks typically accomplished with a mouse or other device in a similar manner to the way one performs pointing in spoken conversations. Such interfaces began with an early system developed by Bolt et al. [1980] in which a large projected screen displayed a number of entities that could be interacted with by pointing at them and providing a verbal command,

such as "Put that there" or "Make an orange triangle". The system allowed for relative commands more akin to natural language, blending the benefits of natural communication with a computer system. For example, referencing prior events in a computer system is usually cumbersome where it is much easier in conversation, so being able to refer to 'the first one I created' for example allows for more flexibility. Likewise, the imprecision of pointing allows for more broad commands that a precise pointing system such as a mouse disallows. The paper specifies commands relating to relative position is possible, for example asking a system to move something to the left or beneath another entity. Given pointing is typically far less precise, this provides an additional versatility to the system.

Such interfaces can also be used to add functionality into the real world. A number of systems exist where the user is able to point of a physical object in view of the system to indicate specified functionality. Early examples of this system had large limitations, such as in [citation needed] where the view did not extend beyond a stereo camera in a small region, meaning only objects in view of the camera that the computer was aware of could be interacted with. Later systems can make use of more complex locational systems and detachable widgets to allow ubiquitous interaction [no reference even found but this must exist]. A computer system can interact with an environment using an interface but this method of interaction removes that layer, allowing for more direct control over behaviour. Applications for this include [citations needed].

This is really lit review stuff... not sure it belongs here.

citation needed, hand in 3D environment

Other such interfaces use pointing in 3D space to refer to objects within a virtual environment, particular 3D environments. presented a system in which interaction with objects was provided by a virtual hand in the computer system, capable of making movements and selections in 3-dimensional space. This provides a solution to the issue of perform 3D selections, given most selection devices only have 2 dimensions of recordable movement, though this style of interface raises an immediate question about the role of pointing in computer interactions.

As previously stated, users will perform a pointing interaction either using a tool to accomplish a task, or as a gesture to further communication. Pointing in a 3D computer environment in most cases operates using similar rules and systems one

uses, such as selection, movement, creation and destruction of entities, but there is no longer a tool and instead a communication mechanism has been appropriated for the task. Most specifically, the use of a physical object which is most associated with interactions with such tools is entirely absent in these systems; and while some require a peripheral such as a dataglove to record the purpose of the interaction, the users hands are free when performing the interaction. In interfaces where operating or manipulating virtual objects is a focus, this form of interaction leans more towards tool manipulation rather than communication.

this probably requires its own, albeit brief section

Tool manipulation in the context of pointing interfaces remains an open question. Without an object of manipulation, the method of indicating the interaction to the computer, as a mouseclick or keyboard command would, now falls to a different modality. Complicating the issue, if the position of a users hand is now used by the system to determine the pointing location, the hand is not free to use another device unless it is attached or carried while performing the pointing. The alternative to this is the performance of a hand gesture small enough to not reduce the accuracy of the pointing.

cumbersome phrase

The section here must lead into the design. Possibilities:

- *Discuss the research questions- lead into experimental? But then how do we section/re-discuss without treading old ground*
- *Consider points a design issue in making such interface (then discuss solution in terms of experiment, or how to solve that problem)*

4.1.2 Design

The first consideration for implementing a pointing interface is whether or not to use a relative or absolute model for pointing:

A relative pointing model is one where the movement of the users hand corresponds to a similar movement of the pointer on the display area. This is the *modus operandi* for the majority of such interfaces due to the simplicity in implementation and the flexibility of the model. A relative pointing system simply has to measure move-

ment, so devices like accelerometers or LED sensors can implement recognition very trivially. The interpretation of movement also allows for factors like pointer ballistics to be used in improving the user experience.

The measurement of relative movement can either be unspecified or relative to a particular position. An unspecified relation of movement is the manner mice and the accelerometers on smart phones work, where the rate of movement and the previous position of the cursor is the only consideration in where the cursor will next appear. This system has the disadvantage that such movement is typically unbounded, so some measure is needed to allow the user to 'reset' the mouse position, or to move the physical device without changing the position of the cursor .

I haven't really discussed cursors, probably for above sections

Some relative pointing interfaces have their cursor placed in relation to another position, which can be seen in devices like the Leap Motion or the Kinect in their typical interfaces. In the former example, the camera captures a specific area that has no relation with the screen itself and recognizes a hand position. The position of the cursor is the hand's position in relation to the sensor position. In the case of the Kinect, using the interface to navigate menus on the Xbox 360, the position of the cursor is the hand's position relative to the rest of the body. This gives the interface additional flexibility, given the hand is the mode used for interaction; it negates the requirement for the user to move to another position to interact with the system, and it allows the user to move freely within the space.

Can we save 'the above'? How do I cross-reference things I've said in my thesis?

An absolute pointing interface can be considered a subset of the a relative interface. Such a system allows users to point in the physical space directly at objects they wish to interact with; if they are indicating towards a shape on a screen, they would point in the direction that shape is in physical space to make reference to it. Unlike relative pointing, which has a layer of interface between the user and interactivity, this model allows users to interact with the system and its icons directly. While this is best known and most successful in the space of touch-screen interaction, we are considering touchless interfaces which are different both in the implementation required and the challenges faced.

This whole paragraph is terrible. Needs to be completely rewritten

Further background, particularly into challenges and issues for absolute pointing interfaces

Is this nec-

This system focuses on the development of an absolute pointing interface. We are interested in this because *The reason we are interested in this, again referencing research question, or some other way?*. In a relative pointing interface, the user has no direct connection with the screen they're working with, as all interaction occurs on a relative plane of motion rather than the one the user physically occupies. The user thus does not interact with the display surface itself, and the relation is

The ultimate goal of the system was to have the system work as an absolute 3D pointing interface; the user should be able to point their hand, arm, finger etc. at the screen and the area they pointed at in physical space should be the area that is then selected, highlighted or has the appropriate operation performed on it.

4.1.3 Implementation Methodology

To begin, our system requires a method of capturing the position of the users arm and body; this allows us to compute the direction of the point. A depth camera or colour camera can be used to determine this information using Computer Vision techniques, and other devices like IR sensors and emitters can work to similar effect, but a wellsupported, accurate and lowcost alternative is the Microsoft Kinect. This device has both a standard colour camera that collects colour data, and a depth camera, as well as an API and set of libraries that convert the video footage into a 'skeleton stream', which identifies any humans in range of the cameras and returns a series of 3D vectors in the camera space that represent the position of all joints their bodies contain.

This is too vague

The kinect is placed next to a display surface with which the interface interacts. We assume for the sake of simplicity in implementation that the camera and the display surface are facing in the same direction, which is typical operating procedure for the device in typical application. We do not need to assume that the kinect is positioned on the same visible plane as the television, but for practical purposes this often produces a more accurate result.

Kinect technical specifications

The technical specifications of the Kinect can be seen in ??

This whole section is too simple and too unjustified

Viable area to flesh out. What are the constructions?

citation needed

With our interface, besides the immediate design constrictions we did not have any strict requirements on implementation. Consequently, an agile development methodology was adopted. Agile development is a system of rapid prototyping, in which a small but achievable goal is stated for a simple prototype, which is then assessed according to the needs of the project to create an incremental goal for the next prototype. This looping phase of assessment and implementation allows a system to match the requirements of the project more closely, especially when the end user requirements of such a system are flexible early in the process.

The process of design followed in this project began with assessing the initial needs of the project. These are outlined in 4.1.3.1 for this initial system. The basis for these constraints was constructed around the first experiment run, which is justified and discussed in ??.

4.1.3.1 Design Constraints

The initial approach to the system had a number of key issues. Firstly, the Kinect has a limited field and range of view, and consequently it cannot perform accurate recognition for any object closer than 80 centimetres from the any sensor. The optimal capture distance for figures is between one and three metres. Having a desktop environment using this system is extremely difficult unless the Kinect is mounted at a further distance, as most users are rarely more than 50cm away from their screen or the capture device in regular computer usage. Consequently, the adaptation of this system for standard desktop use is quite challenging, unless manual figure recognition written for other camera arrangements is implemented.

This doesn't make sense in context

The other major issue the Kinect has with figure recognition, particularly for precision input is the reasonably low resolution of the depth image used for figure recognition. With a field of view of $57^\circ \times 43^\circ$, the depth camera is capable of capturing an area of

To calculate

some amount of metres at 2 metres . . . Consequently, small hand movements in front of a large monitor can have a very low resolution, measurable in pixels. If the display surface has a high resolution, and requires interaction within that resolution.

There's probably a good way of measuring this

At this level, a 100 x 50 button on a 1920 x 1080 monitor exists at a resolution the

camera cannot read.

The simplest solution is to increase the size of the monitor. By using a large wall-mounted display, the area of interaction is larger, which allows the user to make broader movements to reference the entire screen and consequently takes up more space. If the user is sufficiently close to the camera and the screen is large enough to take up most of the camera's field of view, precision can fall into about 2/3rds of the capacity of the Kinect camera, which is high enough for broad pointing.

The later Kinect has improved hardware that mitigates this issue somewhat, but capturing fine hand movements for direct pointing interfaces remains a challenging task for a hands-free interface.

The next limitation is in combining pointing and interacting with this sort of interface. Once an item has been pointed at in an interface, the user will most likely want to perform some sort of interaction with it, which requires an additional input. Given the high precision required in movement to compensate for the low resolution of the screen, it is very challenging to perform recognizable gestures to the Kinect without changing the position of the hand; a common issue in gestural interfaces that integrate both pointing and interaction. Consequently, a way of performing actions without engaging the pointing hand (or body part) directly in the interaction is necessary.

There are several design solutions to this problem. As pointing is typically only performed with one hand in such interfaces (though exceptions exist, see ?, the off-hand is free to perform other actions with the system, either the gestures necessary to interact with the system, or a 'halt' gesture to temporarily stop the pointing recognition so the user can perform other interactions. Another possible solution is to have all gestures be performed solely with the left hand, using a bi-modal interface. While broadly solving the issue, it increases the cognitive load on the user having to do two things at once, and also limits possible interactions to some degree.

Another solution involves the integration of pointing and manipulation into the same gesture. If pointing is possible with a broad set of gestures and a high level of both accuracy and durability can be relied upon from the recognition, the pointing

Tie in with previous calculation when complete

short para

Flesh out, citation needed

Incorrect terminology. Find references to support

hand itself can perform the actual interaction. This has the benefit of being the most accurate metaphor for pointing interaction; typically when a person attempts to access or manipulate an object in real life, the hand that is reaching for that object will also perform the interaction itself, such as picking it up or somehow manipulating it. This would therefore appear to be the most attractive option for our system but it comes with several large caveats. The first is that by having the hand perform the interaction as well as the pointing, it needs to be clear that whatever the hand is doing does not interfere with the pointing as well; outside of very simple finger gestures like gripping and releasing the hand, many gestures require wrist or even arm movement to perform, and this would also change the location of the point by the time the gesture is finally recognized, which leads to usability issues. Even small movements might change the recognition sufficiently, if the target object is small or the recognition is very sensitive. This solution is technically the most desirable but practically would have to be combined with the previous suggestion to be workable.

The solution we eventually went with was the simplest but also the most limited. In addition to pointing with the favoured hand, the participant would hold a physical device like a mouse, a remote-control or something similar. This device should be able to be held comfortably in a number of different hand positions and orientations, and have buttons that can activate certain functions from within the interface. This minimises the challenge associated with performing gestures with both hands and also makes the interactions very easy to perform, as buttons are a very easy vector of interaction. The interaction set is limited to the number of inputs the remote is capable of registering, but it can be extended using chording combinations, Phrasing menus and other features typically seen in pointer-based interfaces. Having an easy-to-manipulate object ensures pointing accuracy is not sacrificed in performing interactions, so it satisfies this limitation sufficiently for our purposes.

This prototype has other benefits for experimental purposes that will be discussed later in the chapter.

The final constraint with our implementation was for the system to have very broad, and customizable recognition. For the system to be appropriate for experimentation, it was necessary for a variety of different interpretations of how to point at the com-

puter screen was necessary. This required the system to calibrate to the user, which required a variety of different methods of pointing to be valid and recognizable with comparable degrees of accuracy. In achieving this, our system would be suitable for observing user variation in pointing gestures.

4.1.3.2 Pointing Recognition

Our system design called for the use of an absolute pointing interface; the position of the pointing is relative to the user's world space rather than the system's interpretation of that space. .

The Kinect skeleton stream is capable of capturing the vector positions relative to the camera, so these allow us to create reference points for a pointing direction. We take two points on the skeleton model and use this to compute a vector that defines the direction the user is pointing. The terminating point of the gesture should be the end of the pointing hand, but the starting point is debatable; the wrist, elbow, shoulder, or even the head can be used appropriately. The most flexible system would make use of the wrist as this allows delicate and ergonomic flexion, but this measure is inappropriate as the hand frequently occludes the wrist when the hand is facing the camera, making tracking and precision difficult. The elbow is very similar to the wrist, but the small wrist movements would only be interpreted as tiny changes rather than large ones.

better define 'world space', and clarify this terminology

Using the head as the starting point is also very attractive as this is how many people point naturally when drawing attention to something specific. In typical non-verbal communication, someone might draw attention to something in the distance or small in the sky by raising their arm straight and pointing direct at it, positioning their head on their shoulder so they are looking exactly where the arm is pointing, so if drawing an imaginary line from head through hand, it would intersect the object. This is the most natural system but it has the disadvantage of being used specifically for this kind of communication ?, where more casual pointing, such as to an area or nearby location, would not use this system. The use case was considered too specific, and thus we did not consider it.

Therebesides, the purpose of this gesture is to ensure the arm itself is the vector, rather than hand to head. For this reason we consider the shoulder to be the most appropriate joint to begin the gesture. It allows both long arm pointing and short arm pointing with reasonable precision, it is close to the eye so the gesture should always be visibly apparent to the user (a feature that using the elbow lacks). A similar approach was used in Fukumoto et al. [1994], and it was the approach we eventually took.

With the localities of the vector specified, the task becomes computing where on a display the user is currently pointing, in the form of a 2D point in screen space. We can easily compute the intersection of the vector with a point on a 3D plane, and the easiest 3D plane to consider and compute is that which the Kinect itself occupies; we define the plane by it's origin (the Kinect's depth sensor) and it's normal (the direction the depth sensor is facing). The Kinect makes use of a lens for it's colour sensor, but the depth values for all joints are adjusted relative to this plane by the Kinect's SDK. This means that while it's not strictly necessary to use this plane for our system, it is convenient for minimizing occlusion, an issue that will be discussed

Is this bad wording?

in more detail later.

Need diagram!

To compute the location of the point, we first define the plane of intersection as

$$\vec{n} \cdot (p - p_0) = 0 \quad (4.1)$$

where \vec{n} is a normal vector representing the facing of the depth camera, \mathbf{p} and \mathbf{p}_0 are two arbitrary points on that plane. For convenience we use the origin of the depth camera for \mathbf{p}_0 .

We next define the point we wish to find, the intersection of the pointing vector \hat{s} with the plane, referred to as s_p . These are defined as:

$$s_p = d\hat{s} + s_e \quad (4.2)$$

where $\hat{s} = \frac{s_h - s_e}{||s_h - s_e||}$ and $d = ||s_e - s_p||$.

The pointing vector (and distance) can be found through algebraic substitution:

$$\vec{n} \cdot (d\hat{s} + s_e - p_0) = 0$$

$$\vec{n} \cdot d\hat{s} \cdot n \cdot (s_e - p_0) = 0$$

$$d = \frac{\vec{n} \cdot (p_0 - s_e)}{\vec{n} \cdot \hat{s}}$$

$$s_p = \frac{\vec{n} \cdot (p_0 - s_e)}{\vec{n} \cdot \hat{s}}$$

This provides us with a world-space point on overall plane where the pointing vector intersects the plane. For a true absolute pointing system however, this is not quite enough to ensure accuracy. Firstly, the display surface must also be on the plane of intersection, so for this reason the kinect should be position on the same plane as the display surface, and facing in as close to the exact same direction as is possible.

Factor in
use of ASM
math; de-
fault math
too messy

diagram?

In the above calculation, we used p_0 to indicate the centre of the plane relative to the Kinect. All values are relative to this central point from the perspective of the Kinect, so if for example typical screen coordinates were being used, the kinect sensor would have to be exactly on the top left corner of the display surface. As well as being a practical impossibility, there are issues with having the Kinect sensor facing directly at the user. Instead, the points must be transposed to within the bounds of a space on the plane that represents the screen's coordinates.

When operating a system, these bounds must be specified by the user through a calibration. A robust calibration process is necessary to account for the need for direct pointing, mobility limitations in the pointing limb and other environmental factors. In the calibration, the user points to each of the four corners of the display surface with their arm. These four points do not form a perfect rectangle in world space in most calibrations; Figures 1 and 2 show two such calibrations, and the shape the pointing produces.

Include
these fig-
ures

As the space that the point falls into isn't reliably mappable with this function, we must translate the position of each individual point from within the calibrated space

to a screen rectangle. Our method of doing so is by defining a screen rectangle, J' with it's position being the extreme-most dimensions of the calibration quadrilateral J :

$$J = J_1 J_2 J_3 J_4 \quad (4.3)$$

$$J' = J'_1 J'_2 J'_3 J'_4 \quad (4.4)$$

Computing each point in J' is done in this fashion:

$$J'_1 = (\min(J_1x, J_2x, J_3x, J_4x), \min(J_1y, J_2y, J_3y, J_4y)) \text{ etc.} \quad (4.5)$$

diagram

This defines J' as containing J , with at least two points in J lying on a side of the rectangle.

For each corner of rectangle J , there exists a vector (for example $\vec{J_2 J'_2}$) that can be used to translate them to their equivalent point in J' .

With this specified, an algorithm is necessary to translate any point Q that lies within J to a corresponding location in J' . Our approach to this is to create a new translate vector, $\vec{Q Q'}$ that is weighted according to the distance between each of the four corners. If, for example, Q is a corner in J , we would simply use the appropriate translation to map it directly to the equivalent corner in J' , but if it lies midway between two points, we would need to use a translation vector equal to the sum of half of the two corner translation vectors to place it in the appropriate position. The distance between the two points determines how much each vector should be weighted in translation.

$$D = \{||\vec{Q J_1}||, ||\vec{Q J_2}||, ||\vec{Q J_3}||, ||\vec{Q J_4}||\} \quad (4.6)$$

$$Sum = \sum_{i=1}^n D_i \quad (4.7)$$

$$\vec{QQ'} = \sum_{i=1}^n \vec{J_iJ'_i} \times (1 - \frac{D_i}{Sum}) \quad (4.8)$$

This calibration process remains imperfect however. It assumes that the quadrilateral traced by users does not have curved sides, which is far less likely given the initial unequal shape of the rectangle is produced by being off-centre to the camera, curvature of arm movements and limitations in mobility when standing stationary. The compromise method was chosen because the consistency of such a shape can't be guaranteed, even with multiple performances, and having too complex calibration process would interfere with usability. Additionally, the position the user is standing in would have a strong impact on the overall shape, so such a process may lead to more inaccuracies than it solves. Conversely, a simpler calibration process would fail to capture the overall shape of the display when traced, which we consider the most crucial aspect of the calibration, so this process is considered a compromise between the two extremes.

We're using a mix of mathematic and computer notation here. I should use one style and stick with it; choose which and sort out how to do it correctly

4.1.3.3 Practical Issues

The Kinect skeletal tracking system assigns a recognized skeleton 3D positions for each joints, along with one of three different states: Tracked, Not Tracked and Inferred. A tracked point is one in which the Kinect can determine where a joint should be on a body, and that point is visible to the camera. This is not always the case; through body movements, various joints or parts of the bodies can be obscured through natural movement, for example standing at an angle to the Kinect hiding one part of the body or having hands move in front of or over each other when gesturing. This is referred to as occlusion of the joint, and without a physical reference point as to where the point is in space, the Kinect must attempt to estimate where it would be. All joints that are occluded, or are otherwise having difficulty being recognized by the Kinect (different body types, unusual posture etc.) change the state of a joint from Tracked to Inferred. While this is not an issues for short periods of time, joints tracked by inference cannot be relied upon for accuracy for any period

of time, so such movements should be avoided during regular usage. For this reason it is recommended users stand at a certain distance from the Kinect so no part of the body is out of the viewing angle of the camera and the body is not so distant as to lower tracking precision, as well as facing the camera forwards and not wearing obstructive clothing.

One issue with using pointing gestures, especially direct pointing gestures with the Kinect is the high incidence of occlusion. If a whole-arm pointing gesture is used, where the arm forms a straight line between the shoulder and hand, the elbow and shoulder both have the potential to be occluded from the camera. As the precision of the shoulder position is necessary to ensure the point direction is correct, this leads to high levels of unreliability in the computed vector, worsening the user experience. Direct pointing is especially susceptible to these imprecisions, as the vector calculation relies upon the difference between the s_h and s_l , which is a very small distance compared the length of \vec{s}_p , so even small variations can lead to high levels of imprecision. Consequently, the Kinect cannot be centrally placed, and instead the final result must be transposed to the display surface.

The last major issue with the system is the resolution of the Kinect Depth Camera. At the highest setting, the IR lens captures images at a resolution of 640×480 which in turn limits how precisely joints can be detected and positioned in space. Making this issue worse is the fact that most of this space is not utilized for an absolute pointing space, instead only a rectangle within that a little larger than the user themselves.

Add figure In Figure 1, a user stands 1.4m away from a camera sensor when calibrating, with each circle representing an individual point on a 17" laptop monitor. Translated to a screen rectangle as above, all pointing interactions are recognized and occur within an 85×50 pixel window of the overall image, which is a very low resolution. When including factors like involuntary muscle movements from the user, the precision of the system is very low.

The only practical solution to this problem is a much larger area for recognition, and to encourage this we used large-screen displays or wall projections to ensure the display is large enough that users are compelled to make larger movements when pointing to each corner during calibration.

4.2 System Performance

This section can give some basic metrics, displays and information about how the system itself is working. This should include: how the system typically functions, constraints on its behaviour, accuracy- a basic Fitts's Law analysis here will provide that information clearly enough. This should serve as an appropriate conclusion.

How do I
bookend
this chap-
ter??

Performance with Pointing

Gestures

bad first title: what is this section even really about?

5.1 Actual and Perceived Interface Functionality

This section is a bridging point to the first experimental design. It explores what is **actually** happening in a system, compared to what the user **thinks** is happening, how this divide is not only necessary but makes for a better system overall. This will need citations to make more substantial, but I don't even know where to look...

- Perception and understanding of gestural system Define 'user understanding' of a system Intent: Explain how a system will have a visible, surface function and an underlying functions

Most interfaces 'abstract' the way they work mechanically with a simpler set of rules for the user to follow - This can be explained through analogy - A typical WIMP mouse interface has a complex underlying model of movement, ballistics that govern what happens when a mouse moves - The user representation is simplified, and the mouse is merely moved in the direction the user wants to click - The more complex interaction changes the 'feel' of the system and impacts the user experience on an intuitive level So the user understanding can be explained as - The explanation the user has been given in how the system works - And their sub-conscious way of interacting with those non-explicit but implemented rules of the system Discuss understanding

in the context of gestural interfaces, and signify their importance Intent: Get across the importance of a user's understanding of the system in ensuring they can use a gestural system both naturally and optimally

What does this understanding mean in gestural systems? - A gestural system has a very complex set of rules underlying how recognition occurs - So system consequently rely having a close correlation between gesture and action to make the system simpler for the user to understand and use - This disconnect, or imprecision in the interaction produces the difficulty in interacting with such systems.

Some one or two papers to discuss how this can be influenced/impacted - Inherent user variation Variance in how a user will see a system in basis of different factors Previous computer experience Previous experience with the specific modality of input Cultural overlap or familiarity with base state - Perception and how users 'view' a system Is there a default state of perception? Can we make assumptions on how a user will initially perceive a system What does this tell us about how the user will use it? How should this information inform the way we design our interface? - Manipulation and modification of that perception Can we change elements of an interface, particularly one that uses free gestures to impact how the user views it? - Findings from Study 1

When designing an interface, be it a device or a piece of software, it is important to understand what the user understands of the system, and how that understanding changes their interaction with it. Relevant to all areas of HCI, this is especially important with gestural interfaces as it directly modifies how a user will interact with a system, and how they may need to limit themselves in order to use it effectively. It is important to examine the issue, which we refer to hereafter as perception.

Consider a hypothetical system that directs a user to restaurants using a map display of their local area. The user specifies what kind of restaurant they wish to visit and the system displays a list of suggestions. In this system, and any such interface be it physical or digital, there are two distinct aspects relevant to how the user operates. First, there is the functionality of the interface, which may include the mechanical and electrical components of the device in use, the algorithms that govern how to

interpret user input, or the logic of the application itself. The second component is the perception, or what the user believes the system to be doing when they interact with it. This can resemble the functionality, but at a higher level, as the detail of operation is rarely relevant to the interface itself. In our example, the functionality of the system may be using locational services to determine where the user is, then performing a keyword search with the provided restaurant kind against a database of listed businesses near to that location, before providing those results to the user ordered by some sort of algorithm, such as proximity, rating, price or previous history. Many of these details are irrelevant to the user's understanding of how the system is working; they only see a map, it asks for kinds of restaurants, and receives a list of appropriate results, without considering how those results came about. The overlap between the functionality and the perception of this interface can be referred to as the 'paradigm' of the system; the obvious rules by which this system works. Except in the simplest of cases however, there are many more hidden rules, that are irrelevant to the paradigm but power the interaction, which we shall refer to as hidden functionality; the user can understand and operate the system with the perception without

The user's perception does not necessarily agree with the user's perception in all interfaces, however, which is a tool used to make user operate an interface intuitively but improving the experience without their knowledge. One example of such an interface is a standard optical mouse. In such a system, the mouse faces a computer, and the paradigm states the user moves the mouse, and they see a proportional amount of movement of the cursor on the screen. How a mouse actually functions is far more complex, with details that aren't relevant, but most interestingly, the back-end implementation of cursor movement is not 1:1, which is what the interface suggests; in fact movement is proportional to acceleration, and other factors that make up the mouse's 'ballistics'. This is because using true 1:1 movement makes the mouse less flexible, and would require either much more movement to move the cursor, or would be less able to do small and precise movements. But a user does not need to understand the particulars of this aspect of the system, and it is much simpler to consider movement of the mouse on a simpler level.

Strong systems are ones where an understanding of the paradigm is sufficient for good interaction, but excellent systems use the paradigm to optimize how a user operates a system, and this is especially true of gestural interfaces. The realities behind the collection of gestures can be quite complex, using databases of pre-recorded gestures for recognition Athitsos et al. [2010], spatial movement ?, or other measures where the instructions relate more to the kinds of movements the user should make for these metrics to operate effectively. . Finding the appropriate balance for the paradigm can be difficult however; presenting gestures too vaguely may lead to users performing interpretations not expected by the system (and leading to worse performance), while too rigid and detailed a paradigm may make the system confusing or cumbersome to operate. In most interfaces, the paradigm can be easily complied with through the addition of interface constraints, that limit certain actions the user can make with the system. The system can ignore certain mouse-clicks or movements, or disallow certain actions from the user. However, in a gestural system, limiting the way a user can physically move is impossible, so having a strong understandable paradigm of interaction that functions effectively with the back-end makes for the strongest system. A system where the paradigm does not correlate well with the back-end means recognition will be low, and one where the paradigm is too close to the back-end will be too complex for basic usage.

More citations, better examples of ambiguity

Our understanding of a system on an instinctive level gives rise to the contentious term 'naturalness' of an interface. A natural interface, we defined, is one that does not require instruction to understand and operate, one that is innate or instinctive to operate. Computers were, at the onset, entirely artificial and unique systems that required extensive training and knowledge to understand how to operate, but as computers went from being a professional tool to a personal tool, the consideration of how systems are presented to users required interaction paradigms that are easier to grasp. The eventual result were a series of interface metaphors, or interfaces designed to resemble things users already had a strong understanding of; most famously the desktop metaphor which treated the file system of a computer like a virtual desk, complete with folders containing files. These changes are still not enough for a user with no computer experience to be able to innately understand how an interface

works, especially if they are also unfamiliar with the metaphor being employed, but it lessens the slope of the learning curve in acquiring the training to use such a system.

Today however, we are able to be confronted with completely new interfaces and be able to operate them with little-to-no training due to our experience. A typical computer user in the developed world likely makes use of their computer with reasonable frequency as learning to use such systems is becoming an increasingly important aspect of daily human life. Where early computer users had to take longer to familiarize themselves with interaction paradigms, the modern user can base their expectations on prior computing experience. This allows systems with well designed and consistent interfaces to have greater expectations on their user to intuitively understand their system, and thus operate it.

I'm sure
I can find
some hard
numbers

5.1.1 Natural Computing

Natural computing however focuses on the operation of computer interfaces without any prior experience or knowledge, and instead relying an intuitive understanding of the interface to discern the interaction paradigm. The use of the label 'natural' is highly contentious, with proponents ? and detractors Churchill [2011] Norman [2010] on it the use of the word. One major argument against the inherent naturalness of gestural interfaces is that due to high levels of cultural variance, gestures humans consider innate vary too widely for any semantics to be consistently applied to them and encompass a consistent understanding among all groups. As discussed in chapter 2, there i a very wide array of gestures that arise in common conversive language, with high varied definitions depending on the group communicating with it, although there exists a small set of gestures that are sufficiently similar across very large groups of humans to be considered to some extent innate. Nevertheless, this small set would not be enough to power a robust gestural system.

flesh out
these ar-
guments.
They're
valuable
and deserve
more than a
footnote

Another criticism comes from the contextual interpretation of a gesture. In addition to gestures having different meanings across groups, they also have different meanings depending on the accompanying speech (the medium through which ges-

tures are most commonly provided). Gestures frequently modulate in definition, and the true purpose of the gesture can rarely be ascertained without other contextual information or accompanying speech, so gesture recognition without these factors inherently weakens the medium.

Further criticisms stems from the application of the word 'natural' itself, and what exactly it means in context. The definition of a natural interface, and what natural means varies from source to source. Many papers define a natural interface as one that mimics or is similar to ordinary human communication, making it much easier to pick up and learn. Wikipedia gives a very different definition, referring to a natural interface as *"an interface that is effectively invisible, and remains invisible as the user continuously learns increasingly complex interactions"*, a description that ties natural interfaces with pervasive computing. In "Brave NUI World", Wixon and Wigdor provide several descriptors of what constitutes a natural interface:

a few citations, they're pretty consistent

cite wikipedia

In the natural user interface, natural refers to the user's behavior and feeling during the experience rather than the interface being the product of some organic process. The production of this conclusion is the end result of rigorous design, leveraging the potential of modern technologies to better mirror human capabilities.

The semantics of the word natural are frequently targeted for criticism. Suggesting computer can be operated naturally is oxymoronic, as computers do not occur in nature. There is no innate way of reacting to a computer without some form of training or way of informing the user as to it's functionality, so it is argued this term does not apply to any set of users.

It can however be argued that interfaces can become 'second nature', one in which there is a minimal cognitive barrier existing between the device being interacted with and the expected result in the system. When a computer system is learned to the desired proficiency, users do not need to think about or consider at a conscious level their method of interaction, and instead can place all focus on the output produced by the system. Such interfaces, rather than being considered natural, instead can be thought of as common, or ubiquitous; apparent, non-natural forms of interaction

that regardless require little attention to operate effectively. Interface devices like mice and keyboard can be put into this category, as they still require practice or even training to use effectively, when this has been achieved, the desired outcome of using the tool becomes the focus of a user's attention, rather than the operation of the tool itself.

In here we have to somehow tie together this admittedly very rambling discussion about what a natural interface is into my own definition. This should relate back to the concept of perception as opposed to function somehow. Read more of "Brave NUI World" to consider this more seriously.

5.2 Study 1

We can assert from this that the following impact serves to create the initial user perception of a natural computing system:

Innate understanding of the system If a natural computing interface is designed to allow any human to understand how it functions, at least on some base level, then a person can be said to have an innate understanding of the system. The impact of this is contentious, as discussed above, and it may be a very small or non-existent aspect in developing a perception of a system.

Environment
issues-right
indent does
not belong
here

Expectations based on experience with tools outside of the realm of computing Although computer systems operate differently to most tools used in other domains, computers nevertheless make use of both interaction metaphors and design aspects from real-world tools to help users bridge their existing knowledge to the unfamiliar domain. On one level, computers employ ways of making their interfaces resemble or mimic a real-world counterpart to provide the user with expectations. One of the most clear examples of this is the desktop metaphor, in which a computer system has a 'desktop' on which all files in the computer can be stored in folders, akin to how one may store real files in an office setting. This concept, rooted in a real-world metaphor is faster to interpret than using a more accurate depiction of a system file pathing system. On another level,

some interface devices are literally designed to resemble real-world devices to make interaction easier, such as pen-mice or light guns that operate similar to those devices, to further help with informing the user. Such measures are useful when the domain knowledge of the user cannot be easily assumed, particularly for audiences where computing system are sufficiently uncommon for this to be a concern (as in populations of very young children or some developing countries).

cite

Expectations based on experience with tools within the realm of computing Although

operating systems, website layouts or interface devices can vary from one computer system to the next, interfaces often attempt to keep functional consistencies to ensure users can apply existing experience seamlessly to their systems. A large number of interface components appear consistently in both functionality and appearance in a broad set of fields. Examples in digital interfaces include interface widgets, like radio buttons, menus and color wheels (two of those examples are also named after real-world systems, again assisting the user in producing an expectation). In physical objects, most computers make use of similar interface devices, and despite some changes, most mice for example function in basically the same manner. Computer usage is sufficiently ubiquitous in most of the world that assumptions on the functionality of commonplace technology can usually be safely made; this has the advantage of not only requiring less explanation of how a system works, but implementations of such widgets are sufficiently commonplace as to make the development much faster than for a unique input system.

Too colloquial

Expectations developed while using the tool Most tools do not simply present them-

selves and leave the user to understand them. Most will convey some information on intended functionality through an instructional process, to varying degrees of thoroughness. This process may be as simple as a written manual of instructions, explaining basic usage of an interface, while others are more thorough, and may involve an interactive tutorial in which the participant operates the tool in a controlled environment with testing. While this is a very commonplace way of introducing a user to a tool's functionality, many interfaces

(especially those branded as natural interfaces) choose to simplify or eschew the process, relying on a simple understanding and allowing the user to fill the gaps in that knowledge with other expectations of function.

What does this tell us about how the user will use it? How should this information inform the way we design our interface? - SEPARATION OF TOOL AND INTERFACE

So we have a perception, a user sees a tool, applies that experience to developing their perceptions and expectations, and they use it. We know both the presentation of the system and, most notably the TOOL provide hints for taht perception. So, ergo, if we change that tool, how does that then impact our perception. And even more importantly, does a change in perception change the usability of the system itself??

5.3 Changing the user perception

With this established, the research question of the thesis can be presented in this light. We can present a hypothetical user with an interface on a computer, as controlled and manipulated by an external tool, which may or may not require them to hold or touch something. The user produces an expectation and perception of the system's functionality based on what they observe of the system itself, and crucially, the tools used to interact with it. If we present the user with an interface that allows the user to operate it according to their perception, and is to some extent flexible to this rule, does the tool used to interact with that particular system change the way the user behaves when operating the system?

This section of the thesis investigates the role tool manipulation has in interacting with natural user interfaces. Hands-free gestural interfaces rely on gestural communication techniques to inform how the system operates, but the provision of a tool to the user can provide the user with additional information and hints on how such an interface should and can be operated. A mouse visibly displayed for use with the interface, for example, suggests that a pointer is the mode of interaction with

the system itself, and users may be confused to be presented with a system where such a thing isn't displayed. Likewise, if the tool presented resembles one which may resemble a tool they are familiar with outside the realm of computing, this may be used to help establish a perception of how that system operates, based on their experience with using that particular tool. For example, a computer system with a controller shaped like a gun may suggest to the user that pointing is used to reference or interact with the system, and the trigger is the method of selection. Conversely, a controller shaped like a musical instrument gives some clues as to not only how the controller operates, but what the purpose of the application is.

Many tools can be adopted to perform a very similar or basically the same task, with variations on how the user operates them. A mouse and a trackball serve an effectively identical purpose but with some changes in the mode of operation, which may impact performance and user preference in varying degrees. When observing this effect with pointing interfaces, many tools can be used for pointing, but the manner in which the tool is used may well change how the system responds to this interaction.

Another important consideration with interfaces that involve large amounts of physical movement are ergonomics. A study looking specifically at the ergonomic considerations of interfaces with and without physical objects Kim et al. [2011] indicated performing a task with a 'virtual object' to be more difficult and fatigue-inducing than an equivalent one performed with a real object, as users. The task of interest required users to move their hand through a physical space to collect a 'cube' and move it to a new position. The participant's hand was obscured with a cover displaying a 3D representation of it in virtual space, as well as the object, with the varying factor being whether the object is real, or is virtually projected. The trial found users extended their fingers earlier and farther with the virtual object despite the representation visible to the user being effectively identical. Where the tool of interaction here changes, in this case from a physical tool to a virtual one, the fatigue induced from operating the system measurably changes. This is a phenomenon that we wished to observe in the context of our system.

Doesn't
bridge into
next section

To better understand the effect the kind of tool has on performing similar or identical

tasks in a pointing interface, an experiment was designed and executed to observe the effect, using the system outlined in Chapter 4.

5.4 A Comparison of Pointing Devices

5.4.1 Aim

We declare the general research question of this experiment as:

"Does the control device presented to the user impact their operation of the interface?"

We define the control device here as being an object the user can hold in their hand, with a series of possible operations such as buttons, wheels or other interface components that have at least some degree of mechanical action associated with the successful manipulation of the device. Interfaces that operate in the absence of such devices we define as such interfaces that capture input about the state of the human body and use this to determine whether or not an interaction has occurred.

Despite the increased amount of technical improvement and commercial interest in pointing and gestural interfaces, both that make use and forego the use of a tool for operation, widespread adoption has yet to occur on the predicted scale. The challenge of making these interfaces increasingly ubiquitous remains formidable, and in the meantime it remains unclear if such interfaces are superior or preferred by users to an object-based interface, in a comparable setting. The ergonomic factors relating to these devices, namely how physically straining or tiresome they are with continued use, remain largely unexplored.

There are three main areas of this question we wished to explore as part of this experiment:

5.4.1.1 Performance

In most cases, performance between interfaces that make use of operating a physical tool and those that do not is varied as the method of interaction itself is typically quite different. As a consequence, the comparison between both methods in a functionally comparable interface is rarely considered, and studies are usually limited to adapting a system to be used in a more traditional interface setting (see Cabral et al. [2005]), or vice versa .

As we are most concerned with the operation of the device specifically, we wanted to observe the performance of a series of different tools performing the same task in roughly the same way, to see if the user's actual understanding of the tool itself changed their interaction with the system. Rather than adapting the interface to fit one way or the other, the tool would instead be adapted to fit with how the pointing works in the system.

Do have
some cita-
tions, find
and list

5.4.1.2 Ergonomics

Research on operating user interfaces requiring the user to perform gestures with arm and finger movements, particularly those in a standing environment, has frequently report strain, fatigue or discomfort in the limbs and other parts of the body as notably impacting the user. The effect has become so prevalent that methods to metricize the issue were reported by HincapiÁl-Ramos et al. [2014], and numerous studies cite discomfort over long periods of use as a major negative factor to system operation. Discomfort can not only hurt the user's preference of a system, it can also adversely affect performance as muscle discomfort can negatively impact pointing accuracy. The experiment was designed to specifically research the discomfort and pain caused as a consequence of operating the devices, to determine whether or the use of a physical object in performing manipulations impacts the strain or discomfort users feel when operating the interface, and if the kind of tool had any impact on mitigating or exacerbating the issue.

5.4.1.3 Preference

Results in the literature suggest the preference for gestural systems can exceed that of their traditional counterparts, for feeling more intuitive to operate, easier to understand or more enjoyable by some other metric. As part of this experiment, we also chose to measure which of our different tools was considered preferable to use and why, to see if any one is substantially considered better by the user, and for what reasons. We hoped to use this information to also give us a better understanding on when to use specific designs of tools in certain circumstances, as research that could be followed up in later experiments.

5.4.2 System

The system used to perform this trial was explained in Chapter 3. We took the technology developed and found several separate devices to operate it based on a series of constraints:

- The device had to have some existing association with gestural pointing
- The devices had to have some distinctions that we hypothesized as changing how the user would perceive them
- A control device was needed that tied the experiment back to a more typical pointing interface, while preserving the system functionality in proper.

At the end of the process, we selected several tools that we felt matched these criteria.

5.4.2.1 Wired Glove

A wired glove is a device that attaches to a user's hand to provide minimal constraint to movement and operation, while capturing hand position, orientation and finger state for interacting with a computer system. These are typically designed as a form of glove worn over the fingers, with each finger of the glove containing a series of bend sensors that measure each fingers' flexion. A glove was selected as our 'hands-free' tool. As a worn device, the hand becomes the tool of interaction, so in a pointing

interface, we suggested users may use their perception of pointing with their hands in other scenarios as how they would approach operating this system.

The Essential Reality P5 Wired Glove was selected for this task. The glove has a large base that mounts over the back of the hand, and five thin rods that attach to each finger by way of several adjustable finger-rings. The glove's bend sensors were used as inputs for the system itself. The base also has a number of infra-red LED's mounted at various positions to inform the system of it's position and orientation by means of a large desk-mounted IR sensor. Although this information was not used in the system, with all positioning information provided by the Kinect, the device required the sensor to be placed reasonably near to the glove for the system to work correctly.

5.4.2.2 Wand

We wanted a pointing device that mimicked a device in real life where the tip of the device would represent the end point of a line used to select or interact with an item. There are many such devices but they often come with additional implications—devices like hose nozzles, guns, pencils were all candidates but they also gave users an impression of the system that may not have been accurate. This is also true for controllers or computer devices that have these perceptions as well. We eventually settled on an imaginary tool that nonetheless we expected people to be reasonably familiar with: a magic wand. A wand is a device that is typically depicted being pointed at a particular thing or place, and something occurring after some activation is performed, like a flick or use of some magic words. Although the 'flick' of the wand wasn't implemented for our system, when presenting it we did so with this nomenclature.

The device we eventually settled on was a custom design of a Nintendo Wii Remote. The Wii Remote is a typical single-handed controller with an array of buttons, but it also features an IR camera that provides the system with position information, although as with the glove this was not used for positioning. The controller made use of a casing to give it the appearance of a "Sonic Screwdriver" from the BBC

Doctor Who series. The device is identical to a standard Wii Remote controller in all other regards, but it has a slimmer, longer and more cylindrical profile. The button inputs on the Wiimote were used for performing tasks, as the use of words would have conflicted with the modalities of the other systems, and a flick gesture may have interfered with the pointing recognition. The appearance of the tool was more akin to a science-fiction device than a more traditional wand, though we maintained this wording for those unfamiliar with the specific prop the device was mimicking.

5.4.2.3 Gyroscopic Mouse

We wished to use a mouse-like device to serve as a comparison point to our experiment. We cannot easily adapt the system to work on a more traditional system, so instead the mouse was adapted to work within the interface. Instead of using an optical or mobile component to measure movement however, a gyroscopic mouse was selected so the mouse could be operated in a 3D environment in a similar manner to the other devices.

The mouse chosen was an OmniMotion Air Mouse, and in addition to having an optical component (unused in this experiment), it contains gyroscopes to measure rotational movement, which corresponds to movement on the X and Y axes of the system. Typical operation is designed to switch between desktop and airmouse mode through use of a button, but as this makes multi-button tasks very cumbersome, this was set to be permanently on. The mouse had to be modified with some considerations to work appropriately with our interface as unlike the absolute pointing interface that powers the other two systems, the gyroscopes provide relative pointing. This included having a 'centering' function that returned the mouse cursor to the middle of the screen.

5.4.3 Method

5.4.3.1 System

The interface was applied to a 3D interactive environment which consisted of a series of tasks that needed to be completed in order. To operate the system, participants began by pointing at each corner of the screen to calibrate the position, followed by completing each of the three tasks in order. The environment was designed to have the appearance of a medieval castle, and the flavour of being involving 'magic' in interactions was continued to further incentivize thinking of the system in terms of direct pointing manipulations. The styling was also designed to engage participants in the experiments and control their perception of the system more rigidly than if the system was more abstract in appearance.

The environment itself used a series of 3D models to resemble chambers and corridors from an ancient castle, where the tasks took place. Although users were interacting with a screen as a pointing device, we wished to determine if users were pointing directly so the cursor was removed from view, meaning users only pointed where they felt they were interacting. Whenever they attempted to select or interact with something on screen however, a small burst of particles appeared around the region that they were pointing at. The color of this burst also indicated the kind of interaction being performed.

There were two interactions available on each device: selecting objects on the screen, and 'grabbing' them to drag to another position. Each was done differently depending on the device; to select with the glove, the user taps their index finger downwards, and to grab the user makes a fist, which they release to let go of the object. With the wand, there is a top-mounted button on the controller pressed with the thumb to perform selections, and a trigger pressed with the index finger that is pulled along with the top button to grab and drag. With the mouse, the left mouse button clicks, while the left and right together drags. The dragging did not require the user to be pointing at the object in order for it to grab, at which point it would follow the mouse cursor until released.

To compensate for the limitations of the kinect's accuracy, rather than being a specific point as with a mouse, each point projected a small collision circle with whatever item was being interacted with.

5.4.3.2 Tasks

In the experiment, participants were asked to take in three trials, one for each device with ordering rotated between participants. Each trial had the same layout and order, and consisted of three section:

NUMBERS
WRONG-
RECHECK
HOW
LONG
TASKS
TOOK

- A series of targets appeared on the screen, of varying sizes and in different positions. Participants selected a total of 30 different targets by pointing at them on screen at this stage. The positions and sizes were always randomized, but the seed was saved as part of the results. In this trial, the targets appeared as ghosts that each selection was dispelling.
- At the bottom of the screen, a key of varying size and position would appear. Two locks on either side of the screen, also of random position but non-random size also appeared. The goal was to select the key with the pointing device, then drag it until it overlapped the lock. Once it overlapped, it disappeared and another key appeared, with the required lock to drag swapping sides. This had to be completed a total of 10 times.
- A cage appeared in the centre of the screen, and around it were a number of fireflies. One of them shimmered a particular color. The goal was to select the shimmering firefly, and then drag and drop it inside the cage. Once it dropped in the right position, it disappeared and another firefly would start to shimmer. Ten different fireflies had to be dragged into the cage. Only six were ever on screen at once to avoid over-cluttering the view.

what was
this about?

Each interface had a slightly different

The trials were run consecutively, with a short test period before the first trial, and a break between the first and second trial. The trials were designed to have the

appearance of being performed in a gothic ch teau, with the tasks themselves resembling ones being accomplished by casting magic spells; for example the select-drag-drop task involved casting a spell (selecting) at a floating firefly, and dragging him with another spell to land inside a cage. This allowed users to treat each device as operating according to their own expectations, and adopt an approach that felt contextually natural in operating the devices.

Corresponding to the length of each given task to keep them fairly consistent, a total of nine selection tasks, six select-drag tasks and five select-drag-drop tasks were performed for each device. All tasks for a single device were presented to a user consecutively, with a brief break between each task explaining what needed to be done before it was begun. At the end of each sequence of tasks, the participant was given a 5-10 minute break to rest their arm before continuing with the next interface device.

At the beginning of each trial, the user was required to calibrate the device, and re-calibrations were permitted in the event the device performed contrary to user expectations or with otherwise imprecisely or inaccurately, as was occasionally the result of a bad calibration. Before the selection trial and the select-drag trial, the user was given the chance to practice pointing at the screen and performing selections and grabs in the absence of a selectable or draggable object, to ensure the device was performing correctly and to gain familiarity with how the interface worked. 15-45 seconds were allocated to these pre-training sessions, during which the user was requested to move the cursor to various parts of the screen and perform each motion several times, after which the trial would start. In the event issues arose during any of the tasks that hindered performance, the trial was reset and users were given several minutes to rest before starting again.

During the task completion phase of the experiments, the mouse cursor is hidden from participants. Selection gestures are accompanied with a visual indication on screen in the form of a small burst of stars in the area of the selection, so users know precisely where their gesture was directed and by how much to correct it, if necessary. The grabbing gesture can also be used to display the cursor position, but without a selected object this has no effect on the trials.

5.4.3.3 Measurements

During each trial, the system keeps track of every selection and grab attempt performed and the position of the cursor at each attempt, along with a time signature. Each screen state was also rendered to an image file, which included size, orientation and position of the targets in each scene. The footage of each experiment was captured by the Kinect for each section of the trials for every participant. This data included skeletal positions in Kinect space and depth images of the participants interacting, which could be used for later analysis. At the end of each section of the trial, the user was also queried on the discomfort or pain felt in their arms, and where specifically the discomfort was located, and how severe it was. A Borg CR10 scale was used in questioning on the amount of exertion they feel, where 0 represents minimal or no exertion and 10 represents maximum exertion and discomfort.

no cite yet

A questionnaire was also provided to participants for qualitative measurements, specifically their preference in pointing devices. Basic demographic information was recorded to begin. At the end of each section, the user was also asked to fill out a short section about the device they had just used, including how intuitive, natural and physically exhausting it was to use. These were rated on a Likert scale from 1 to 7. Finally, a series of short answer questions were presented at the end of the experiment, where devices were compared and general comments regarding the interface were made. The complete questionnaire can be seen in Appendix B.

5.4.4 Results

The first run of this experiment was run with a total of 12 participants. 8 male, and 4 female. All were between the ages of 21 and 36, with a mean age of 25. All participants were computer science students, and reported regular computer use. Most had some familiarity with the Kinect, with 5 participants reporting high familiarity with the device and 4 participants reporting limited experience. Other gesture capture technologies such as the Nintendo Wii were also familiar to participants, with only 1 participant having never used any other motion capturing devices. None of the participants reported existing back or shoulder injuries that may influence the

	<i>Glove</i>		<i>Mouse</i>		<i>Wand</i>	
	Time	Errors	Time	Errors	Time	Errors
μ	136.8	50.3	149.0	71.5	79.7	45.8
σ	62.1	18.4	81.3	33.8	33.8	20.7

Table 5.1: Mean and standard Deviations for time and errors in completing selection tasks for each device

fatigue of the trial, though one suffered discomfort from continued standing. 11 of the 12 participants were right handed, and one reported as being left-dominant but relatively ambidextrous.

5.4.4.1 Performance

The pointing section of each trial (performing the initial selection) was designed to allow for an analysis of the index of difficulty for each pointing device using Fitts's Law . However, the decision to make cursors invisible along with smoothing added to the pointing recognition to reduce inaccuracies from the Kinect led to the results showing no strong trend for most participants. As a consequence, the performance relied heavily on the accuracy of the Kinect and personal skill with the interface, rather than visual feedback, which made such an analysis less meaningful. Another metric recorded was the number of incorrect selections and overall time for tasks. Both of these metrics showed stronger trends and better correlated with questionnaire results and fatigue, so these values are used as a measure of how accurate and easily each participant was able to use each device. Certain outliers were pruned where appropriate, such as participants having to recalibrate the Kinect, or taking time to rest if they were suffering from fatigue.

citation
missing

Table shows the means and standard deviations of each metric over every selection task performed by the users. Error rates between the glove interface and wand were roughly equivalent, but users on average took substantially longer in performing the operations with the glove. The mouse was on average the poorest performing device, but also the device with the greatest level of variance. Figure 1 shows the comparison of total errors over total time taken for each participant, showing the general spread of performance. The wand can be seen here as having generally the

X

best and most consistent performance, though all devices are demonstrated as being able to perform well, with relatively few errors and very quick; the fastest participant completed the total of 21 selection tasks with the wand in 29.4 seconds. Two outliers from mouse usage were also pruned; Despite being informed that accuracy was a metric of success, some participants chose to rapidly click their devices to keep the cursor visible when making selections, which made their data unusable.

Later drafts will need to include figures. These have been excluded in this draft for simplicity.

5.4.4.2 Ergonomics

During the trial, a variety of postures and stances were adopted. Users stood and were typically quite still, with some shuffling. Likewise, movement in the spine and back was relatively uncommon, and most motions were made with the pointing arm. The method of control and orientation of the arm was consistent enough between participants to be separated into three categories, ordered from lowest reported induced fatigue to the highest:

- The shoulder at rest, with the elbow kept at the side and the forearm extended on the transverse plane, pointing at the display. Holding the device in this manner, users made a combination of wrist and small forearm movements to control the pointing. When holding a device this way, users experienced very little fatigue, usually located in the biceps.
- The shoulder partially at rest, the elbow extended from the side, with the forearm raised and the wrist bent to face the device to the display. Forearm rotation and wrist movements were the primary method of interaction. Extended use of devices in this orientation was associated with moderate fatigue in the biceps and triceps.
- The arm fully extended with the elbow locked. The user moves the entire arm at the shoulder in this position. High levels of fatigue were associated with this orientation, typically in the triceps, shoulder and neck muscles.

Would a diagram of the three planes of movement be a useful addition around here, or further back?

While users typically kept a consistent arm orientation for the trials, it was not un-

	<i>Immediate Fatigue</i>	<i>Continuous Fatigue</i>
Non-parametric Test	$\chi^2 = 7.023, p = 0.03$	$\chi^2 = 11.286, p = 0.004$
Pariwise Glove vs. Wand	$p = 0.008$	$p = 0.015$
Pairwise Glove vs. Mouse	$p = 0.04$	$p = 0.005$
Pairwise Wand vs. Mouse	$p = 0.55$	$p = 0.232$

Table 5.2: Dstribution-free analysis of questionnaire results pertaining to immediate and contiually-induced fatigue from each device

common for a change in posture. This was most commonly moving from an outstretched arm to a relaxed or partially relaxed arm, likely a measure to combat fatigue, or from a relaxed or partially relaxed arm to an outstretched arm, usually due to issues with accuracy. One or two participants would also support their right hand with their left at certain points during the trial to combat the fatigue. Figure 2 shows the prevalence of various orientations with each device, by counting the number of times and the duration of the trial during which the stance appeared for all participants.

On average, the glove interface was reported to be the most fatiguing by participants ($\mu = 4, \sigma = 1.9$), followed by the wand ($\mu = 2.4, \sigma = 2.2$) with the mouse being the least stressful to use ($\mu = 1.3, \sigma = 2.5$). The mouse results are slightly misleading; the mode for the device is 0 with 8 occurrences, but one participant reported a rating of 9 for discomfort due to holding the device in an unusual orientation- this was located in the fingers, not in the arm or shoulder muscles.

Immediate fatigue induced from use, and the extent of continued fatigue from use for each device was also reported in the questionnaire. The results are shown in Table 2, with no significant difference between the wand and mouse ($p > 0.05$), but more pronounced variation between the glove and all other devices. This confirms the glove as being the most fatigue-inducing device to operate.

In the questionnaire, users were asked to grade each interface system on a seven-level Likert scale, how natural, intuitive, learnable, reactive, accurate and generally easy to use the interface was perceived to be. Of those, only two yielded a statistically significant result in non-parametric analysis. In both instances, pair-wise analysis revealed no significant difference between the glove and mouse. The wand however

Just use another graph type then!
Enough with these silly averages

was found to be more intuitive than the glove ($p = 0.016$) and the mouse ($p = 0.018$). The wand was also found to be easier and quicker to learn than the glove ($p = 0.024$). The fact that all other results lacked a clear correlation suggested personal preference played a large part in these reports.

Users were also asked to rank each interface in terms of preference, and provide comments as to those rankings. The results of this are shown in Figure 3, indicating a general preference to the wand interface. Comments for preference to the wand referenced its appearance or comfortable grip or being easier to press buttons than on the mouse or with finger gestures with the glove. When the mouse was preferred over the wand, it was typically due to being the less strenuous to operate, with complaints of the mouse being its unpredictable behavior and difficulty in reaching buttons. Preferences to the glove were reported as being the more natural of the interfaces. Complaints with the glove included a sense of being unresponsive and fatigue-inducing. All devices received comments regarding their accuracy, either being higher or lower than other devices- this appeared to be a highly subjective measure and there was no agreement between participants on this.

5.4.5 Discussion of Results

Our expectations were that users would find differences in induced fatigue and preference between the devices, but pronounced differences in performance and speed were not expected. The results suggest the presence of an object held in the users hand has a profound impact on how users manipulate the system, and the amount of effort users expended. The study also showed remarkable levels of variation in user preference.

The results from observations suggest that interfaces with a physical object held in hand tended to encourage users to use a less direct form of pointing, performing points with the upper arm fully or partially at rest, and the elbow at the side, while having no object encouraged users to keep their arms fully extended towards the display to attempt to ensure better accuracy – yet their results are no better. The wand was observed to also be subject to in-hand manipulation; users were observed

turning or rotating the device within the palm of their hand rather than performing the equivalent rotation with their wrist or forearm. One participant went so far as to keep the elbow almost totally still while turning the wand as much as 45 degrees in the palm of the hand- the slight wrist movements that accompanied these movements were sufficient to inform the interface of the pointing direction. Overall, less motion was observed in video feedback with using the wand when compared to the glove interface.

This rather counter-intuitively seems to be accompanied with a better performance index for the wand. Less physical movement by the user is required overall to move the pointer to the desired area, but as the region of selection is far smaller, inaccuracy was expected to be higher. This may be attributed to the fact that selection and grabbing was reported to be much easier with the wand and mouse when compared to the glove. Users on average seemed to exert more energy attempting to perform selections, as the glove required a strong, deliberate motion to select an object on screen compared to the other two devices. Some users commented that the act of performing the taps for selections or the fist clench for grabs would lead to involuntary arm movement that hindered accuracy. The sharpness of the gesture was intended to avoid undesired select gestures being performed but this remains an issue of concern for hands-free interfaces with binary-state interactions. That a strong correlation can be seen between increasingly extended arm, trial length and fatigue matches expectations, and remains an important consideration in interface design.

5.4.5.1 Trial Caveats

The poor results for the mouse can largely be attributed to the absence of a visible cursor. As the mouse tracked relative movement rather than absolute movement, it could not be discerned without performing a select or grab to show where the cursor was, so overshooting targets and erratic motions appeared with many participants. It was also common to click extremely quickly to constantly display the position of the cursor. The rotational nature of the mouse sensors also seemed to confuse participants, many attempting to move the mouse with lateral motions that

had unpredictable results on the cursor position.

The variation in user preference reflects the diverse background of participants. Comments such as “familiarity with the Wii may have been useful” appeared in the questionnaire, despite the fact the system shared no similarity in functional operation to the Wii, but the method in which users operated the interface reflected that experience. Preferences for any device appear most heavily influenced by the devices that users found to perform most effectively, though there were some instances where this was supplanted with the justification of “familiarity” with the existing device, particularly with the mouse.

5.5 Followup Studies

This previous study found that depending on the device being held in performing pointing gestures (regardless of how the gesture is being captured by the system), users will perform the gesture with a different posture and motion. The wand, which could be used as a more direct form of pointing device encouraged users to rely on the device to perform the pointing, and in so doing did not physically engage themselves as much with the task as a device that either failed to provide this function, in the previous case being the wired glove, or provided pointing through a separate paradigm in which such physical activation is not typically necessary, as was the case in the mouse interface. This engagement from our observations manifested in more deliberate and less ambiguous pointing by way of increased flexion in the forearm and upper arm from a rest position, which demonstrates an increased reliance on the user’s own anatomy in ensuring precision and accuracy in operating the interface.

Devices that the user holds that already resemble pointing devices, in this case a Wiimote encourage the user to adopt a more at rest stance and use the length of the controller to aid in the gesture. Conversely a device that does not perform this role and has no ability to be used directly in pointing, in the previous experiment a P5 Wired Glove, the user will rely on their arm and hand more heavily on ensuring the

gesture is precise and on target. This occurred regardless of the form of tracking being employed; even though both approaches were tracked using camera and the users arm rather than the device itself, this trend still appeared. As a consequence, the level of fatigue participants experienced when operating the different interfaces and the speed and accuracy with which they could perform tasks was also impacted.

Provided evidence that different aspects of the user's actions and behaviour appeared to be influenced by the environment in which the controller was presented, the goal of our research became finding specific aspects of that environment that we could manipulate and observing the effects more thoroughly. In the previous experiment we had explicitly been modifying the kind of tool the user had been operating, but we now sought to check our features of the environment.

- Controlling the posture of the user, by requiring them to either be seated or standing and seeing how this impacts the movement of their arms
- Controlling the position of the fingers on the hand, specifically when pointing with a bare hand or non-constricting device. Determining if having the index finger outstretched in the typical posture for full-arm pointing changes the way the user observes the devices.
- The use of different devices, specifically one that best fits the environment we presented it in.

In exploring each aspect, we maintained the same style of interface and all other aspects of the experiment remained the same as in section 5.3. By coming to a better understanding of how we can impact a user's behaviour, how their fatigue can be changed by the presentation of the interface, we could understand how to make such pointing faces as usable as possible.

5.5.1 Changes in Study Method

Several changes had to be made to the design of the experimental methodology, in improving on the original experimental design. There were a number of recognition issues with the software, in which certain body types were difficult to recognize with

the kinect and the results became less accurate with continued use of the device.

The biggest change made was changing the origin of the ray cast from the user's point (to determine where the point is being performed) from the elbow to the shoulder. We found that particularly when users had a fully outstretched hand, the elbow was often occluded, and the precision of the pointing was lower with the elbow than it was with the shoulder. We also found that the size of the screen being operated was still quite small, and to make the system more usable on the screen, we artificially increased the size of the pointing area by a small amount. This moved the corners of the screen out from the bounds of the view area calibrated, so perfect absolute pointing was not retained, but the movement was still equivalent, as was the directionality with the only difference being slightly more movement was necessary from the user for the interface to respond as desired. The benefit of this was a broader range of movement allowed for higher precision in trials that needed it, particularly the drag-drop trial in which the targets were very small.

Several other small changes to the software were made to increase usability. These included slightly increasing the smoothing applied to movement, in an effort to further mitigate the imprecision from the device, a selection system to change which arm is used in pointing for left-handed participants, and changes to the color of some objects to make the software more accessible for visually impaired participants.

One of the major issues with the previous trial was the proper usage of the wired glove. The glove has adjustable finger rings to change where the glove grips the hand and make it more accessible to users with different sized hands, but for those with slim fingers and smaller hands, the rings often fit loosely, so manipulating the flex sensors was much more difficult, requiring a very firm grip. It was also very difficult to find an appropriate threshold for activation of the flexion sensors useful to all participants. If the sensors operated with too high a threshold, users had to form a very firm grip to activate them, which was fatiguing when done repeatedly. Conversely, having a lower threshold of activation led to the sensors remaining active even after the user released the grip, forcing them to extend their fingers back to deactivate the sensors, which also caused discomfort and was difficult for those with small fingers. Due to these issues, several participants were not able to complete

the experiment. It was therefore necessary to find a more accessible and usable replacement for this device, that still allowed for minimal device intrusion and an analogue to hand-directed pointing without the caveats of this previous device.

The replacement ultimately chosen was a Neo Reflection Wireless 3D Finger Mouse, pictured in the figure below. The mouse attaches to the index finger through a plastic ring, and is manipulated through gyroscopic sensors inside the mouse that measure pitch and yaw, as well as two touch-sensitive buttons that can be pressed with the thumb. The device has a very small form factor of 58cm x 22cm x 13cm, and weighs 13 grams, so it does not greatly intrude on the hand's movement. The buttons are used in place of the finger bending for the glove, which allows for a substantially less exertion in operating the interface. The device also provides several different sized finger rings to ensure different hand sizes can make use of it without looseness or discomfort. Although the device is quite noticable when working, nevertheless it was seen as a reasonable alternative to a non-intrusive worn device for pointing with the hand rather than a device.

Add picture

How to
bookend
this section?

5.5.2 Standing and seated postures

In the previous experiment, we had found that users suffered considerable discomfort from the way they held the device up. With many of the interface devices, they had to keep their arm fully extended, which caused large amounts of arm fatigue. If making use of a tool or device that encourages this sort of interaction, having a method of discouraging it could make the interface substantially more usable. Our first approach was to change the user's posture altogether, in a method that actively encouraged a more relaxed method of interaction without hampering the system usability. The most immediate and feasible change meeting those criteria is changing whether the user is sitting or standing.

For this experiment, the participants ran through two separate trials; in one they would be standing up when operating the interface, and in the other they would be seated. The order of these tasks was switched for each user. A chair was selected

I need a
small sec-
tion here
discussing
the er-
gonomic
benefits of
being seated
over stand-

that had no arm rests, so that the movement of the arms wasn't obstructed.

In this experiment we were looking specifically at fatigue and discomfort in operation, so we adopted two metrics; at the beginning and end of each trial we queried the user about their perceived fatigue, and took the difference as the total fatigue either accumulated or relieved over the course of the experiment. We were also specifically interested in how the user shaped their arm when pointing at the screen so we developed an additional metric, based on that shape. The user's skeletal position data was collected by the Kinect, which included their forearm, elbow and shoulder joints. The amount of in the elbow, measured as the angle, informs us of what sort of arm shape has been adopted in operating the interface. A small angle suggests the 'swan neck' posture, with a partially extended arm and the hand pointing towards the screen. A larger angle suggests further stretching of the forearm out towards the screen, with a complete 180 angle having the arm fully outstretched in that highly uncomfortable posture experienced by users of the glove in the previous experiment. This angle was captured while the users were operating the interface over the course of interaction and included in a rolling average. Periods of rest or stopping were flagged in the footage, and excluded from this average. This provides us with a closer metric, on average, to how the user behaved with their arm over the experiment, rather than simply logging examples of where we saw that behaviour exist.

This sub-experiment was run with a total of 6 participants. The group was entirely male, between 18 and 29 years of age with a median age of 23. All members of the group reported regular computer use, and two reported being experienced with gesture-based interfaces.

5.5.2.1 Results

The mean angle of the elbow have been captured and recorded and then normalized between the smallest and largest angles recorded by that participant.. We report the standing and seated average angle, as well as the difference- a positive difference indicates the angle became larger in the seated trial, and a negative difference reports

For absolutely no reason- find the old data and fix this

<i>Participant</i>	<i>Standing</i>	<i>Seated</i>	<i>Difference</i>
1	0.830808	0.766563	-0.06425
2	0.822841	0.637888	-0.18495
3	0.711392	0.697496	-0.0139
4	0.647317	0.76847	0.121154
5	0.752027	0.595249	-0.15678
6	0.583414	0.498825	-0.08459

Table 5.3: Difference in normalized arm angle between standing and sitting trials

<i>Participant</i>	<i>Standing Fatigue Increase (per minute)</i>	<i>Seated Fatigue Increase (per minute)</i>
1	0.25316	-0.857143
2	0	0
3	0.73171	-2.089552
4	0.27273	-0.727273
5	-0.872727	0
6	1.49254	-1.096606

Table 5.4: Difference in normalized arm angle between standing and sitting trials

a smaller angle.

In measuring fatigue, the experienced fatigue was captured at four different time points; one before and one after each of the two trials. As before, a short break was given to users between each experiment to ensure fatigue did not modify the other results. Fatigue was reported on a Likert scale, between 0 and 7. The difference in these two reports was then compared to the amount of time the user spent performing the test to determine the experienced fatigue increase over each minute of the trial

5.5.2.2 Discussion

For the sitting versus standing experiment, we found that standing participants experienced more fatigue than when seated. We found that the mean angle of the elbow in pointing increased on sitting and arm was more extended. This is attributed to issues with the encumbrance of the posture. A higher incidence of occlusion of the elbow and shoulder was observed during the trials, as the seated posture allowed it

Again this is a terrible metric, and assumes a consistent change in fatigue. Find the results and replace with a simple fatigue gain over the experiment, and a separate how

to be tucked behind the body more easily. This is what we attribute to the smaller angle, and the comparatively straighter arms used when gesturing from a seated position. However, this correlated with an overall lower amount of reported fatigue than standing, despite finding a straighter arm correlated to higher incidence of fatigue in our previous experiment.

This actively suggests that changes to the user's posture will impact their general performance with gestural interfaces, and that having users be seated rather than standing not only decreases the amount of discomfort they have from extended standing periods, but also encourages them to use their arms in a more ergonomically appropriate way, which in turn allows for extended use.

Pretty short conclusion, but I don't know how much more there is to say?

cite

5.5.3 Hand shape

When considering pointing as a gesture, it is most commonly thought of as an outstretched arm with the hand posed to draw attention in a specific direction or manner. The gesture, as discussed in chapter 2 is perhaps best known and most familiar to us when the hand forms a fist and the index finger is outstretched; this is so familiar and powerful a gesture that in some cultures it is considered rude when directed at other people. When asking a user in an experiment to, for example, point at something in the room without any devices present, this is the gesture most expected to be seen. However, as found in our earlier experiment, this gesture is most associated with a fully outstretched arm, which in turn we associated with a larger amount of induced fatigue.

The use of a device that frees the hand to perform this sort of pointing was associated with assuming an extended arm position. When the user is holding a device, the device can be used to perform the pointing, while the arm remains at rest. By suggesting the gesture to the user through our devices, we were able to change their behaviour in some ways, so further exploration of this concept, particularly in the context of index finger pointing, was proposed.

In this experiment, participants ran through two separate trials. In the first, they would be equipped with the finger mouse and asked to perform the trial as usual. In

<i>Participant</i>	<i>Without Splint</i>	<i>With Splint</i>	<i>Difference</i>
1	0.479929	0.585284	0.105354
2	0.882051	0.860367	-0.02168
3	0.618639	0.657432	0.038793
4	0.898231	0.884259	-0.01397
5	0.89681	0.859148	-0.03766
6	0.530339	0.61404	0.083701

Table 5.5: Difference in normalized arm angles between free hand and splinted hand

the second, a small wooden splint was attached using pull ties to both the tip of the index finger, and the base of the finger. This splint was designed to keep the index finger straight during the experiment, so the hand worked as though it performed a traditional ‘index finger point’

We were most interested in the posture adopted in the arm of participants during this experiment, so we used the same metric taken in the standing and sitting sub-experiment. Fatigue was also measured by the same methodology.

5.5.3.1 Results

The metrics used to measure arm angles are the same as in the previous sub-experiment. The angle of the elbow was measured and then normalized according to both the smallest and largest angle the participant used in that experiment.

5.5.3.2 Discussion

The results of this experiment did not produce a pronounced effect on participants. There was no observable trend between participants that suggested a link or effect with this technique. We cannot therefore make any conclusive statement about directly controlling the gestures a participant makes use of.

This trial had a very specific modification that may have failed to show results for

Add some qualitative measures here on how much people did that particular gesture

I need AC-TUAL analysis of this. Maybe Tom can help me come up with some

a number of reasons. Index—finger deictic pointing is highly ubiquitous in modern culture, but it is not universally recognized or used as a pointing gesture (true for at least one cultural group of indigenous Australians, see Wilkins [2003]), although it is highly unlikely any members of our trial would not be familiar with the gesture in this context. The other possibility is the contextual disassociation of the gesture; we hypothesized that by having the hand in this orientation, much like the devices in the first experiment that encouraged a more relaxed arm posture as the device seemed to imply the pointing behaviour that the system made use of. If this was not occurring in our experiment, it may be that as nothing else about the trial changed, the splint was dismissed as irrelevant to the system's function and ignored in further performances, leading to a null result. Most participants showed a very consistent arm angle through the experiments, which suggests they were comfortable with their method of using the system.

5.5.4 Identifying Optimal Attributes for 3D Interfaces

The work in previous sections has identified certain changes in behaviour and operation that make certain interfaces change how users perceive them over others.

We've also seen some evidence that certain features of these interfaces is desirable for our technology in different contexts—our devices are different, and they are used differently so what makes them best suited for a given interaction task?

Most importantly for interacting with (in this case) 3D interfaces, what features are the best suited for the job? What are least well suited?

5.5.4.1 Introduction

The 3D environments and data visualizations have become a common element of modern computer usage. From their basic and most common setting in video games, to virtual interactive spaces such as second life, and even in more contemporary settings such as search engines and map displays, the need for users to be able to interact with data in 3D is now commonplace. This has allowed a more intimate

way of being able to view and manipulate data, as well as providing a more engaging interface for the user. With this, however, has come an increased interest in finding novel and better interface methods to interact with this data; as typical inputs are constrained to 2 dimensions, additional input is required to perform navigation and movement in the additional dimension presented. Interfaces which better relate to the environments the user is interacting with are of chief interest among these. A large area of research arising from this are ubiquitous, gestural interfaces; users operate no device but use a series of sensors to capture information about their movement, interpreted as interaction with the system. However, setup and use of equipment necessary for these sorts of interfaces is still quite high for the typical end user; though exceptions exist, device manipulation remains the typical method of interacting, even in 3D spaces. There exists a class of devices that can be used in a standard 2D setting, but are operated by the user in a 3 dimensional space. These devices also use a suite of sensors to capture human movement but with the device in-hand, users are easier able to make precise interactions such as button presses for interacting with and manipulating data, as well as avoiding the ambiguity that comes from classifying and working with gestures, while still retaining the richness of 3D interactivity. In this paper we will be exploring the aptitude these devices have for operating in 3D environments, and what elements make them suitable for this sort of operation. Several devices have been selected with various elements that we hypothesize are important in working in 3D environments for analysis and comparison with a user base to determine the most important considerations in selecting appropriate devices for interfaces, or in designing such devices material.

5.5.4.2 Methodology

The research question of this paper is, *when using a device to perform pointing an object manipulation in a 3D environment, what factors of the device impact its optimality?* We define a device as any non-ubiquitous object manipulated by the user to interact with the system, and its optimality by the factor of how well it performs, how consistent it is found to perform and to what extent it is favoured or preferred by users. We set out to explore this problem with an experiment, measuring optimality of a series of

interfaces in an attempt to isolate these important factors.

System The environment the experiment has been conducted in is presented to the user as a bounded 3D space. Within this space, a series of objects are presented to the user which they are able to interact with, given a set of rules explained at the top of the screen and by the experimenter. These include the ability to select specific objects on the screen, and to pick up and drop them, locking them to the on-screen cursor until the user lets them go. The graphical style of the interface is that of a castle, with large rooms connected by corridors. The environment is not navigable by the user; instead the system determines where the user is to be taken depending on what stage in the trial they have achieved. The environment was displayed on a wall-mounted flat screen television, with users standing 1-2 meters away from it. A cursor was displayed to indicate where the user's pointer was positioned as a rotating star. The star would release a burst of smaller stars when a selection action was performed by the user, and would change color from green to red when the user was attempting to perform a drag. Objects on screen would also change colour or otherwise make it clear if they were being selected and dragged successfully.

There were a total of three variations on the tasks that users were asked to perform:

- Selection tasks, in which the user was asked to select a series of ghosts that appeared in random positions and sizes on the screen. Selecting each target would cause it to disappear and a new one to appear. Ten ghosts needed to be selected to progress.
- Grabbing tasks, in which the user had to select a key appearing at the bottom of the screen and drag it over a target area (a locked door). The key had a consistent size, so this exercise judged moving the object around in 3D space.
- Dropping tasks required the user to select a series of randomly positioned and sized fireflies in a room, grab them and drop them inside a small cage at the centre of the screen.

The selection interaction was required for each experiment and was used as a base measure of performance in the trials. The drags and drops, being more difficult tasks

to perform were measured separately.

Input Devices Three separate input devices were tested in this experiment. Figure 1 shows each of the three devices, as well as the hand dynamometer used in analysing arm strength and fatigue.

Gyroscopic Air Mouse The Omni Motion Air Mouse can operate as both an optical and gyroscopic pointing device. It is roughly the same size and button layout as a standard desktop mouse. An internal gyroscope is capable of detecting sensitive rotation of the device, and this was used exclusively by participants in moving the cursor on the screen with this device. The device was held in the palm of the user's hand, with the thumb being used to access buttons. The left and right mouse buttons were used for selections and grabs respectively.

In our experiment, the gyroscopic mouse is the closest analogue to a traditional mouse used in typical computer usage, and was the one expected to be most familiar to users.

Wired Glove and Camera A combination of an Essential Reality P5 Wired Glove and the Microsoft Kinect were used to produce a direct pointing interface in the simulation. The Kinect sensor captured the skeleton position of the user, and used the vector between the elbow and wrist to project a line that determined the direction the user was pointing. By calibrating the users pointing to a monitor surface, the user was able to point their arm from any orientation and any position to move their cursor on the screen.

The actual selection and grab interactions were captured by the P5 wired glove, worn on the right hand. A series of bend sensors along the fingers were used to capture the flexion of each digit. In our experiment, flexing the index finger was used to indicate a selection, while forming a fist and flexing all digits performed a grab. As the glove is strictly handed, only right handed or ambidextrous users were able to participate.

This device provided the greatest degree of user-system interaction, as pointing was direct and user-defined, so besides explaining its operation, almost no

training was necessary to learn to use it.

Finger Mouse The Neo Reflections Wireless 3D Finger Mouse operates with the same technology as the gyroscopic mouse, i.e. an internal gyroscope that measures rotation to move the cursor. However, the device is substantially smaller in size and is mounted by a small plastic mount to the index finger of the user. Only small movements of the index finger are necessary to move the cursor rather than full hand movements. The device is also substantially lighter, and remains dormant unless the user touches their thumb to the outside of the casing.

Being the lightest and most sensitive of the devices, we anticipated this would have very high performance. Of interest to the outcome of this experiment is that both this device and the Gyroscopic Air Mouse can also operate as typical optical mice.

5.5.4.3 Experimental Methodology

Each experiment conducted had a participant use two of the three devices in sequence. Each trial involved ten selection tasks, ten grabbing tasks and ten grabbing and dropping tasks. After completing each task in a trial for a single device, the participant was permitted a 5 minute break before performing the second.

At the end of each experiment, users were asked to fill out a short questionnaire providing their preference for each device on a Likert scale and in a series of qualitative comments. In addition to this, on the suggestion of an occupational therapist, a grip dynamometer test was employed at the end of each trial to measure changes in arm strength over the course of the experiment (in the event fatigue physically weakened participants).

In addition to this, the program recorded all actions the user took in operating the system. As targets and objects displayed in the visualization were in 3D, target size and distance from cursor were calculated in post-processing of data. This information informed us of how quickly users made selections and grabs, how long each

<i>Device1</i>	μ Time (s)	σ
Air Mouse	2.68	1.54
Finger Mouse	2.84	1.18
Glove	3.30	1.75

Table 5.6: Mean Times and Standard Deviation for each Device in Performing Selection Trials

task took and how many errors or mistakes were made in the process.

5.5.4.4 Results

A total of 18 participants were used in this experiment. Each couplet of trials had a total of 6 participants, so each device was used and reported on 12 times.

Performance Each device was analyzed according to a Fitts’s Law projection, but none produced a sufficient correlation consistently for their values to be meaningful in this experiment. This was expected to be for two reasons: firstly, the difficulty of tasks never had a great variance in difficulty, or $(\log 2(1 + \frac{D}{W}))$ as measured by the law, as very small targets were considered unfeasible for selection in the interface. Secondly, a degree of smoothing was applied to the ballistics of cursor movement; this was added as a necessity to the Glove interface due to imprecisions in the Kinect and carried over to other devices to ensure consistency. The addition of this smoothing we believe encouraged users to focus on precision, despite being informed the trial was timed; this shows as error rates are very low for all devices but a much higher variance in time taken. Analysis of devices using Fitts’s Law, therefore, was unhelpful.

Instead, we compared the overall time taken for each device in the trials to determine which tended to perform the fastest. We also used the standard deviation to determine how consistent each device was to use, and whether or not there was a great deal of variation in its movement. Tables 1 and 2 show the results of this analysis.

For the first set of trials, performing basic selections of objects, the fastest device was found to be the Air Mouse, with the Glove having the worst performance. However,

<i>Device2</i>	μ Time (s)	σ
Air Mouse	2.85	1.88
Finger Mouse	4.03	3.29
Glove	4.81	4.21

Table 5.7: Mean Times and Standard Deviation for Each Device in Performing Grab and Drop Trials

while the Finger Mouse was notably slower, it had substantially less variation in user performance, suggesting it could be used more consistently in these trials.

The second set of trials was less ambiguous, with the Air Mouse having both superior speed and consistency in its results when compared to the other devices. The trials also show the glove to be the worst performing interface. This could partly be the fault of the bend sensors on the glove, which were at times unresponsive (particularly if participants had small hands), and also an issue with the Kinect which occasionally suffered from a lack of precision that at times made tasks (especially the grab and drop tasks) very difficult for users. Several instances also existed in trials of users having to stop and rest their arms or reposition themselves in an attempt to reset the sensor.

Preference and Fatigue In the questionnaire, users were asked to rank their preference of the two interfaces they were given, and list reasons why, as well as comment on any fatigue they may have experienced and issues they had using any of the interfaces. Of the three input methods, the Air Mouse was the most popular; Eight of the eighteen participants stated it was their preferred device. Six of the eighteen participants indicated they preferred the finger mouse, while two stated the glove was their preferred interface. In giving feedback on issues or reasons for their preference, we were able to identify 5 general areas participants discussed in their feedback:

- Imprecision with the interface, too sensitive, slow or unresponsive
- Fatiguing to use and unnatural feeling.

Figure 2 shows how many times each device was mentioned in user comments with these issues.

<i>Device</i>	<i>Change in Reading</i>
Air Mouse	1.34
Finger Mouse	2.10
Glove	0.07

Table 5.8: Mean change of average dynamometer recordings from baseline to after trial with device

<i>Device</i>	<i>Mean Fatigue</i>
Air Mouse	2.66
Finger Mouse	4.00
Glove	5.92

Table 5.9: Reported fatigue from each interface on a scale from 1 to 7

The most commonly reported problems were to do with precision and sensitivity of the interfaces, having difficulty in selecting objects with ease or speed. Other reports included problems with the equipment failing to recognize movements; this was an issue solely of the glove that had set thresholds for accepting bends that users had difficulty in determining. Fatigue was reported by some users but only mentioned as a usability issue in the instance of the glove. Interfaces were also referred to as 'unnatural' or 'strange feeling' at times, and this seemed to be a subjective measure.

Dynamometer data was taken from participants but results were not consistent enough to be reportable. Participants would both increase and decrease from their baseline after trials, by varying amounts. As a result, we found the dynamometer was not a useful measure of fatigue or effectiveness after using these devices. The change in readings from baseline dynamometer results after every trial are listed in Table 3.

Users did however report fatigue for each device on a scale in the questionnaires. The mean fatigue participants experienced with each device is listed in Table 4.

Discussion In observing the performance of each device it is clear the gyroscopic mice greatly outperformed the camera-based interface. While some of this may be due to imprecision in the bend sensors on the glove, it is more greatly attributed to the far higher level of precision the mouse was able to produce over the Kinect. Issues with skeletal tracking with the Kinect camera would often lead to the pointing location changing erratically during trials, as the elbow and shoulder are frequently

eclipsed from view when participants operate the interface. Measures such as changing posture and position came some way to fixing this but the interface was never as smooth and precise as the gyroscopes. It was surprising to see a substantially higher performance with the air mouse over the finger mouse, given both operate with effectively the same internal technology. Part of this success may be attributed to the device being larger and thus easier for participants to handle, or the analogue between it and using a standard desktop mouse (which many commented as being one of its strengths in the questionnaires). Differences in simple selection are relatively small, but in performing the more delicate task of moving objects around the screen, the mouse was found more precise by participants.

The differences between the results between the Air and Finger mice is interesting, since we expected better results with the Finger mouse, as the weight of the device is reduced, and so user pointing would be more similar to natural pointing with our hands. This contrary result is consistent with results in a paper on wands and other holdable devices [6]. It seems likely that holding a device is beneficial. In our future work we can investigate device size and weight and time trade-offs, as surely there are a maximum sizes/weights that meaningfully affect the way users point, and length of experiment, which could objectively demonstrate fatigue effects of different devices.

The results of the questionnaire were similarly revealing to the priority of the user when considering the design principles of the experiment. The majority of criticisms levelled at the system were to do with difficulty in performing small or precise movements with the system; where these interfaces acknowledge and work with the fact they suffer from poorer precision in these environments but provide a greater level of immersion and connection to the interface as a tradeoff. With only 3 comments regarding this (and no positive comments), this suggests that users prioritize performance, particularly precision over the user experience.

Performing with Freeform Gestures

Summary your thesis and discuss what you are going to do in the future in Section 6.1.

6.1 Future Work

Good luck.

Conclusion

As of yet unspecified content, I guess..

7.1 Who even knows...

Bibliography

- ARGYLE, M., 1988. *Bodily Communication*. Methuen & Co, New York, 2 edn. (cited on pages 6 and 7)
- ATHITSOS, V.; WANG, H.; AND STEFAN, A., 2010. A database-based framework for gesture recognition. *Personal Ubiquitous Comput.*, 14, 6 (2010), 511–526. doi:10.1007/s00779-009-0276-x. (cited on page 34)
- BOLT, R. A., 1980. "put-that-there": Voice and gesture at the graphics interface. *SIG-GRAPH Comput. Graph.*, 14, 3 (1980), 262–270. doi:10.1145/965105.807503. (cited on pages 12 and 13)
- CABRAL, M. C.; MORIMOTO, C. H.; AND ZUFFO, M. K., 2005. On the usability of gesture interfaces in virtual reality environments. doi:10.1145/1111360.1111370. (cited on page 42)
- CASELL, J., 1998. *A Framework For Gesture Generation and Interpretation*, book section 11, 191–215. (cited on page 6)
- CHURCHILL, E. F., 2011. Missing the point. *interactions*, 18, 5 (2011), 80–83. doi:10.1145/2008176.2008194. (cited on page 35)
- CĂLDRAȘ, C. AND SHAH, M., 1995. Motion-based recognition: A survey. *Image and Vision Computing*, 13, 2 (1995), 129–155. (cited on page 13)
- FUKUMOTO, M.; SUENAGA, Y.; AND MASE, K., 1994. "finger-pointer": Pointing interface by image processing. *Computers & Graphics*, 18, 5 (1994), 633–642. doi:10.1016/0097-8493(94)90157-0. <http://www.sciencedirect.com/science/article/pii/0097849394901570>. (cited on page 25)
- GOODALL, J. v. L., 1968. The behaviour of free-living chimpanzees in the gombe stream reserve. *Animal Behaviour Monographs*, 1 (1968), 161–311. (cited on page 7)

-
- HEWES, G. W., 1976. The current state of gestural theory and language origin. *Annals of the New York Academy of Sciences*, 280, 1 (1976), 482–504. doi:10.1111/j.1749-6632.1976.tb25512.x. <http://dx.doi.org/10.1111/j.1749-6632.1976.tb25512.x>. (cited on page 8)
- HINCAPIÁL-RAMOS, J. D.; GUO, X.; MOGHADASIAN, P.; AND IRANI, P., 2014. Consumed endurance: a metric to quantify arm fatigue of mid-air interactions. doi:10.1145/2556288.2557130. (cited on page 42)
- KIM, Y.; LEE, G. A.; JO, D.; YANG, U.; KIM, G.; AND PARK, J., 2011. Analysis on virtual interaction-induced fatigue and difficulty in manipulation for interactive 3d gaming console. In *Consumer Electronics (ICCE), 2011 IEEE International Conference on*, 269–270. doi:10.1109/icce.2011.5722577. (cited on page 40)
- KNAPP, M. L. AND HALL, J. A., 2010. *Nonverbal Communication in Human Interaction*. Wadsworth Cengage Learning, Boston, 7 edn. (cited on pages 6 and 8)
- LICSÁR, A.; SZIRÁNYI, T.; KOVÁCS, L.; AND PATAKI, B., 2006. Tillarom: an ajax based folk song search and retrieval system with gesture interface based on kodály hand. doi:10.1145/1178745.1178760. (cited on page 12)
- NORMAN, D. A., 2010. Natural user interfaces are not natural. *Interactions*, 17, 3 (2010), 6–10. doi:10.1145/1744161.1744163. (cited on page 35)
- WILKINS, D., 2003. *Why Pointing With the Index Finger Is Not a Universal (in Sociocultural and Semiotic Terms)*, book section 8, 171–214. Psychology Press. ISBN 0805840141. (cited on page 63)
- WIXON, D. AND WIGDOR, D., 2010. *Brave NUI World*. 1 edn. (cited on page 11)
- WOODWARD, J. C., 1976. Signs of change: Historical variation in american sign language. *Sign Language Studies*, 10, 1 (1976). (cited on page 12)